

FinInteract: The Single-Gold Illusion and the Elicitation–Integration Gap in Ambiguous Question Answering

Anonymous submission

Abstract

Large language model agents are increasingly deployed for financial information retrieval and hence downstream analytics or investment decisions. But financial terminology is specific, and a question without highly specification may admit multiple defensible readings with materially different values, e.g. Meta Platforms’ “operating income” is \$46.75B consolidated or \$62.87B for the Family of Apps segment, but each reading’s ground truth is exact and programmatically verifiable. A capable agent should recognize the ambiguity and ask a clarification question instead of committing to a plausible yet unintended reading. However, existing financial benchmarks generally attach only one gold answer per question, failing to distinguish between agents which resolve the ambiguity from those making ungrounded guesses, which we call blind spot *single-gold illusion*. To address this illusion, we release FININTERACT, a bilingual (English/Chinese) benchmark of 173 instances that pairs each question with a *default* and an *intended* interpretation (fixed by a disambiguating context) across a five-category ambiguity taxonomy, and grades whether an agent elicits and then integrates the right clarification. Across a range of proprietary and open models, we observe that re-grading the *same* outputs against the intended answer leads to accuracy drop by more than 3 times, verifying the single-gold illusion. Our fine-grained analysis shows that this failure is interaction as models actually answer with over 90% accuracy once the interpretation is resolved but only about 45% under their own clarification, suggesting that the key bottleneck lies on correct clarification and downstream integration. We further include skill metric that localizes failure modes of different financial ambiguities across two well-powered categories (entity scope and metric definition) and three exploratory ones.

1 Introduction

Large language model (LLM) agents are increasingly utilized in financial analysis to answer questions in regulatory filings (Islam et al. 2023; Chen et al. 2021; Reddy et al. 2024), carry out multi-step research and retrieval over corporate disclosures (Bigeard et al. 2025; Choi et al. 2025; Krumdick et al. 2024; Xie et al. 2024). The agent reads a

user’s question, searches through relevant filings and returns the answer for downstream actions.

However, financial questions are often deceptively underspecified. Financial terminology is highly specific, and a natural language question may map to several defensible readings for different numbers, e.g. in the question “*What was Meta Platforms’ operating income?*” (Figure 1), an agent must decide whether the user means the consolidated GAAP figure (\$46.75B for FY2023) or the Family of Apps segment (\$62.87B), both being programmatically verifiable against SEC facts. A capable agent should recognize the ambiguity and follow up with the user instead of directly committing into a plausible but unintended reading.

Existing financial benchmarks cannot measure this as they grade only the final answer by assuming the question is disambiguated (Islam et al. 2023; Chen et al. 2021; Reddy et al. 2024; Krumdick et al. 2024; Xie et al. 2024). Without grounding check, they cannot distinguish whether the agent actually resolves the ambiguity or guesses the common reading. We call such blind spot *single-gold illusion* (Table 13), which is currently underexplored in the financial domain relative to general ambiguous QA literature (Min et al. 2020; Li et al. 2025; Kuhn, Gal, and Farquhar 2022). We study financial ambiguity over three research questions:

- **RQ1:** Do frontier models recognize financial report ambiguity?
- **RQ2:** What financial ambiguity leads to most failures?
- **RQ3:** Can we incorporate interventions at inference time to improve resolution?

Answering these questions requires three design properties. First, the benchmark must draw representative ambiguity from real filings. Second, it must expose fine-grained dimensions of ambiguity through paired default and intended readings. Third, benchmark construction should have low cost and is safe from potential leakage to keep answer out of disambiguating context, with verifiable ground truth validated by domain experts. (Sec. 2).

We release FININTERACT, a bilingual benchmark in English and Chinese containing 173 instances, classified into five financial ambiguity categories. Each question Q is paired with a *default* interpretation recording a possible non-expert assumption and an *intended* interpretation fixed by the disambiguating context. By grading the same model outputs

This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.




Figure 1: **Evaluating financial agents on ambiguous queries.** (1–2) A real query (“What was Meta Platforms’ operating income?”) looks fully specified but hides ambiguity across five categories (entity, metric, period, filing, policy), here yielding two defensible answers: a default \$62.87B (Family of Apps segment) and the intended \$46.75B (consolidated GAAP). (3) A search-only agent returns the plausible default, which a single-gold benchmark scores as correct, the *single-gold illusion* that overstates competence by up to $3\times$. (4) FININTERACT instead grades whether the agent *asks*: the GPT-5 agent reaches 20.2% under free interaction versus 93–95% once the ambiguity is resolved, so the gap lies in interaction, not missing knowledge.

against the default and intended interpretations, we effectively expose this single-gold illusion. FININTERACT includes skill metric to localize which clarification ability the model lacks, and evaluates the agents under a ReAct protocol where the user simulator answers only from the disambiguating context (Figure 2) to measure the agent’s quality of clarify-then-integrate action.

Our evaluation answers the research questions in turn. First, while frontier models detect ambiguity, they rarely resolve it effectively, as they ask the correct clarification question but still **failing to convert into the correct answer at the end** (Sec. 4.4). Second, **policy-related ambiguities** are the least resolved, as none of the proprietary models would initiate clarifications when facing this sparse category. Third, **inference intervention facilitates ambiguity resolution**, increasing accuracy from about 44% to above 90%, suggesting that this benchmark is not only diagnostic but also actionable. (Appendix C). This paper makes three contributions

1. **Problem definition.** We formalize ambiguity in financial questions with five categories grounded in standard accounting and disclosure practice (Sec. 2).
2. **Benchmark.** We introduce FININTERACT, a bilingual financial ambiguity dataset containing 173 instances, each paired with a default and an intended interpretation verified by domain experts. (Sec. 3).
3. **Empirical findings.** Our evaluation shows that existing models commonly fail to return correct final answer even when they have correctly clarified the intention, which can be improved via intervention at inference time. (Sec. 4).

2 Problem Formulation

Given an ambiguous financial question instance (example in Table 1), it should contain the ambiguous question over a company filing Q , the disambiguating context C which is the minimal set of scope or definitional constraints that collapses Q into the unique intended answer A under intended interpretation $\mathcal{I}_{\text{int}}$. The ambiguous question typically contains correct answer $A_d \neq A$ under the default interpretation $\mathcal{I}_{\text{def}}$ which a non-expert would assume when reading Q . Specifically, $\mathcal{I}_{\text{int}}, \mathcal{I}_{\text{def}}$ are structured records specifying the financial ambiguity taxonomy defined below.

Table 1: This example of ambiguous financial question contains three plausible entity scopes, i.e. consolidated total and two reportable segments, leading to an entropy of 1.58 bits. Without the disambiguating context C , a model’s default interpretation leads to the wrong answer.

Field	Value
Q	What was Meta Platforms’ operating income?
C	Total company (consolidated) results under U.S. GAAP for the year ended December 31, 2023.
A (intended)	\$46.75B [<i>consolidated GAAP</i>]
A_d (default)	\$62.87B [<i>Family of Apps segment</i>]
Ambiguity category	Entity scope
H_0	1.58 bits
Intended span	“Income from operations for 2023 was \$46.75 billion. . .”
Default span	“Family of Apps. . . Income from operations \$62,871. . .”

Financial ambiguity taxonomy. We derive five categories of financial ambiguity from standard accounting and disclosure practice (Table 2). Each instance should normally exercise one primary category with one or two optional secondary categories in case of additional ambiguity.

Table 2: Five-category financial ambiguity taxonomy.

Category	Definition
Temporal scope	Ambiguity over which time period is intended: fiscal year ending September vs. calendar year; Q3 results vs. full year; TTM vs. point-in-time.
Metric definition	Ambiguity over accounting basis: GAAP net income vs. non-GAAP adjusted income; organic revenue vs. as-reported; operating EBITDA vs. net income.
Entity scope	Ambiguity over consolidation level: segment vs. consolidated total; parent-attributable vs. full consolidated; Class A vs. Class B shares; A-shares vs. H-shares.
Filing vintage	Ambiguity over filing version: original 10-K vs. amended 10-K/A; as-reported vs. restated; preliminary earnings release vs. final filing.
Recognition policy	Ambiguity over accounting policy: revenue recognized point-in-time vs. over time; gross vs. net revenue; capitalized vs. expensed R&D.

Disambiguation entropy. We quantify how ambiguous an instance is using the prior confidence entropy $H_0 = \log_2 k$, where k denotes the number of plausible interpretations of Q without clarifying interaction, e.g. the example in Table 1 has three plausible interpretations, hence contributing to an entropy of $H_0 = \log_2 3 \approx 1.58$ bits. This further serves as our metric to evaluate difficulty and efficiency in Sec. 3.5. While both A and A_d are defensible from publicly available filings, an agent should be able to identify the ambiguity by reading Q alone without C , where this pairing between default and intended interpretation generalizes the target-distractor design from INTERACTCOMP (Deng et al. 2025) to the financial domain.

Agent Interaction Model. We follow INTERACTCOMP’s ReAct framework (Yao et al. 2023) where the agent may issue three types of actions, namely **search**(Q) to retrieve a passage from the filing corpus, **interact**(Q) to present a yes-or-no question to the simulated user who answers based on C , and **answer**(Q) to submit a final answer. Specifically, we evaluate agents in three primary modes (Table 3) and four ablation modes. We additionally include a category-conditioned interaction mode, described in Appendix C.7.

3 Methodology

FININTERACT extends formulation in Sec. 2 as a semi-automated ambiguous financial QA corpus built from public filings. Each carries a programmatically verifiable answer, a paired default and intended interpretation produced by an

Table 3: We evaluate interaction modes including answer-only (A), answer with search (A+S) and answer with search and interaction (A+S+I).

Mode	Description
A	Pure parametric recall
A+S	Oracle retrieval augmented
A+S+I	Standard full interaction
Always-ask	≥ 1 follow-up question
Category-oracle	Provide primary category
Template-oracle	Fixed human-written question
Enumerate	Lists all interpretations and answers

LLM construction pipeline with minimal human annotation and validated by domain experts. Figure 2 shows a worked instance and how the benchmark scores an agent’s two resolution paths.

3.1 Design Goals and Principles

Design goals. FININTERACT targets four goals: **(G1) Representative coverage** to draw real ambiguity from public filings. **(G2) Fine-grained diagnosis** to contain diverse possible readings so that evaluation can localize which clarification ability a model lacks. **(G3) Low-cost, leak-safe construction** to build and verify the corpus at scale in a semi-automated setting. **(G4) High, human-validated quality** to validate the instances and judges by domain experts with minimal annotation cost.

Design principles. To satisfy the four goals, we build FININTERACT under four principles: **1): Easy to verify, Ambiguous to resolve** to create programmatically verifiable ground truth using SEC XBRL facts and CNINFO-structured data but the question cannot be correctly resolved without disambiguating context C , **2): No yes/no questions** so that the agent has to deliver a grounded reasoning support, **3): Includes company name** to ensure no ambiguity in entity, i.e. capitalized proper noun for English and CJK company name for Chinese, and **4): Answer is not in disambiguating context C** so the agent cannot take shortcut by copying directly from the context.

3.2 Data Sources

We build the instances from official filings. For English instances which take about 80% of the passage pool, we build them from SEC filings, DocFinQA long-context 10-K/10-Q passages with verified QA pairs and XBRL-derived facts for S&P 500 and Russell 1000 companies within the financial year of 2024 to 2025. For Chinese instances, we build them from A-share annual reports of SSE 50 and CSI 3000 constituents using CNINFO and akshare structured APIS. We avoid reusing published QA pairs directly to reduce memorization risk, though we do not claim absolute contamination safety. The current evaluation corpus contains 173 instances over 61 companies within 2022 to 2025 filings, labelled with ambiguous taxonomy from Sec. 2, with detail breakdown in Table 4. Specific details regarding passage counts and difficulty distribution are included in Appendix H.




Figure 2: **FININTERACT** rewards actions that resolve the intended interpretation. (1) Each instance carries a primary ambiguity category and two defensible readings, a default (Family of Apps, \$62.87B) and an intended (consolidated GAAP, \$46.75B), separated by a hidden disambiguating context C . (2) The agent either answers without clarifying (Part A, a plausible but wrong reading) or asks a question that a user simulator answers only from C (Part B, correct elicitation). (3) **FININTERACT** scores three signals: final-answer correctness against the intended interpretation, interaction quality (did it ask, target the right category, and resolve intent), and the single-gold failure of rewarding the default.

Table 4: FinInteract dataset statistics. Disambiguation entropy $H_0 = \log_2 k$ over k plausible interpretations.

Property	Value	Property	Value
Total instances	173	Category: entity	94
Language: EN	53	metric	63
ZH	120	recognition	9
Source: EDGAR	50	temporal	7
CSRC-CNINFO	120	filing	0
DocFinQA	3	# cat.: single	128
Filing: 10-K	53	two	45
annual report	120	H_0 : mean	2.20
Filing yr: 2022	40	median	2.00
2023	53	range	1.0–3.6
2024	58	Blind-solve rate	0.7%
2025	19	Unique companies	61

3.3 Curation

We curate each instance from a three-role pipeline, where a **constructor** agent using GPT-5 generates the full (Q, C, A) triple together with the default and intended interpretations and category label from a filing passage. Specifically, we include category specific instructions to prevent drifting towards categories that are easy to generate. Next, we conduct a **sanity check** on code level to reject malformed passages instantly without API calls. Lastly, we employ an **adversarial verifier** to run multiple trial rounds to answer the question Q

without disambiguating context C and reject instance where more than two trials produce the intended answer and interpretation simultaneously. Details regarding the prompt, sanity checking codes, output schema and cost are included in Table 15, Appendix H and Appendix G.

3.4 Validation

We employ finance-literate annotators, including Chartered Financial Analyst (CFA) charterholders and licensees from Hong Kong Securities and Futures Commission (SFC), to review instances to check whether 1) question Q is ambiguous and 2) disambiguating context C uniquely identifies the intended reading without stating the answer and 3) both intended and default answers are correct. We record a substantial agreement with an average 0.85 Gwet’s AC1 (Gwet 2008), indicating strong chance-corrected consensus that, unlike Cohen’s κ , stays reliable under the heavy skew toward “yes” labels. The full protocol and annotation details are in Appendix H.

3.5 Metrics

We include standard metrics from ambiguous QA and interactive retrieval literature (Chen et al. 2021; Deng et al. 2025; Islam et al. 2023; Wei et al. 2025) to directly compare **FININTERACT** to prior benchmarks.

Accuracy (%) measures the fraction of instances whose final answer matches the *intended* interpretation, graded by

GPT-4o-mini with financial domain tolerance including measurement difference, entity equivalence across different stock markets, current normalization or fiscal year matching.

Default-capture rate (%) is the fraction whose final answer instead matches the *default* interpretation, which measures how often a model returns the naive reading so we can reconstruct what a conventional single-gold benchmark would have scored (Sec. 4.3).

Interaction diagnostics includes the interaction rate (IR) which measures proportion of instances with more than one clarification question, average turns of clarification (Round), and accuracy on predicting the true ambiguity category (AC@1). Specifically, we report AC@1 by whether the question names any true category and whether it can name the single most important category.

Confidence measures the Expected Calibration Error (ECE) between model confidence and correctness to capture agent’s overconfidence when reading the ambiguous queries.

Full metric definitions, including the accuracy grading tolerance, the AC@1 decomposition, and a supplementary clarification-efficiency score (DisE⁺), are in Appendix A.

4 Experiments

4.1 Evaluation Setup

Agent architecture. We evaluate agents under the ReAct (Yao et al. 2023) framework under seven configurations from Table 3. The **interaction** action presents a yes-or-no question to the simulated user where the user has to respond “yes”, “no” or “I don’t know”.

Prompting strategy. FININTERACT evaluates interactive agents under different interaction protocol, including answer-only, answer with search, answer with search and interact, and the category conditioned policies. All models run under the same ReAct system prompt. Additional information is in Appendix C.7.

Models. We evaluate over a diverse models such as proprietary models including GPT family with GPT-5.6-sol, 5, 4o and 5-mini, Claude-Somnet-5, Gemini-3.5-flash, Grok-4.5, DeepSeek-V4-flash, GLM-5.2 and Kimi-K3, and open source models including Qwen3 family including 4B, 8B, 14B, 32B and two mixture-of-experts models namely Qwen3-30B-A3B and Qwen3.5-35B-A3B. We evaluate the proprietary model via OpenRouter.

4.2 Overall Performance

Table 6 shows that even frontier models are vulnerable to financial report ambiguity. Specifically, we see that no models exceeds 44% accuracy on the intended interpretation under free interaction mode, consistent across proprietary and open source models. The opportunity to interact does not help to improve accuracy, as the interaction difference Δ is non-positive for most models, indicating that existing models do not fail in asking correctly, but in converting the clarified intention into the correct answer. This is supported by the 95% interaction rates and high AC@1 indicating that the models do recognize what kind of ambiguity it is to ask the correct clarification question. Specifically, we conduct further ablation studies in the interaction by forcing multiple interaction

Table 5: Forcing more clarification does not help with accuracy, but *resolving* the ambiguity lifts every model to 93-95%. This gap suggests that it is the model’s inability to elicit *C*’s itself but not missing knowledge or grounding (Same abbreviation as Table 3).

Model	A	A+I	Forced(4)	Oracle
GPT-5	0.6	20.2	1.5	95.4
GPT-4o	0.6	4.6	0.0	93.1
GPT-5-mini	0.6	0.0	0.6	95.4
Qwen3-30B	1.2	11.6	–	90.2
Qwen3.5-35B	0.6	28.9	–	91.9

rounds and directly providing the disambiguating context *C* as oracle, Table 5 shows that forcing more rounds of clarification actually lowers the model accuracy, which may imply agents getting lost in multi-turn conversations (Laban et al. 2026). Moreover, under the oracle scenario, the models can answer at about 93 - 95% accuracy consistently. As the context *C* does not contain final answer *A*, the improvement suggests that the gap is not missing knowledge, but the ability to disambiguate appropriately for downstream reasoning. We include the full finding in Appendix B. **Finding 1: Models are capable of observing the ambiguity and answering correctly in oracle-clear intention, but fails to guide itself during intermediate conversion between updated intention and correct answer.**

4.3 Failure of Existing Benchmarks

By re-grading models on default readings from existing benchmark and the intended reading FININTERACT curates and verifies, we observe that the accuracy differs significantly (Table 6) with a consistent inflation on the default scenario. This suggests that if existing benchmarks are curated without specific linguistic updates, model performance in financial QA might be conflated. We include details of the findings along with human-annotated study in Appendix B. **Finding 2: Conventional single-gold financial QA benchmarks fail because building benchmarks with only one gold answer would likely encode default reading but not human-preferred intentions.**

4.4 Systematic failure from recognition policy

We extend the study into each ambiguity group. Table 6 shows that in the high AC@1 shown in Table 6, while models are capable of determining ambiguity category in most cases, it consistently fails in the recognition policy group. This motivates us to study this blind spot of existing models by examining the following case study. Given the question “What were General Electric’s revenues?”, the answer is either \$67.95B or \$26.79B under different recognition policies, the former being GAAP revenue that combines over-time service recognition with point-in-time product recognition, and the latter counting only revenue recognized at a point in time. Models fail here as they consistently return a single figure without asking for the expected recognition basis, defaulting to a generic assumption that “revenue” is well defined

Table 6: **(a)** Main results, reporting accuracy (%) against the *intended* interpretation per mode, with Δ the +Interact minus +Search difference (full-scale rows are $N=173$ and the cross-vendor frontier rows, IR and ECE blank, are a preliminary $N=50$ pilot). **(b)** Re-grading the same outputs against the intended versus the default reading a conventional benchmark would treat as gold (Same abbreviation as Table 3). **(c)** Per-category targeting (AC@1), where recognition policy is a shared blind spot (= 0 for every model). Category sizes: entity (Ent.) 94, metric (Met.) 63, recognition (Rec.) 9, temporal (Tmp.) 7.

(a) Main results							
Model	Acc (%)		Answer+Search+Interact				
	A	A+S	A+S	Δ	IR	AC@1	ECE
<i>Proprietary models</i>							
GPT-5.6-sol	10.0	36.0	30.0	-6	—	.98	—
GPT-5	0.6	6.4	20.2	+13.8	98.8	.86	.53
GPT-4o	0.6	12.1	4.6	-7.5	99.4	.94	.82
GPT-5-mini	0.6	0.0	0.0	0.0	96.0	.94	.00
Claude-Sonnet-5	4.0	46.0	44.0	-2	—	.98	—
Grok-4.5	2.0	40.0	32.0	-8	—	.87	—
GLM-5.2	12.0	40.0	40.0	0	—	.96	—
Kimi-K3	16.0	38.0	34.0	-4	—	.98	—
DeepSeek-V4-flash	12.0	28.0	24.0	-4	—	.96	—
Gemini-3.5-flash	14.0	6.0	30.0	+24	—	.96	—
<i>Open source models</i>							
Qwen3-4B	0.0	8.1	12.1	+4.0	28	.85	—
Qwen3-8B	0.6	14.4	24.9	+10.5	15	.88	—
Qwen3-14B	0.0	18.5	22.5	+4.0	36	.74	—
Qwen3-32B	1.2	27.2	8.4	-18.8	61	.96	—
Qwen3-30B	1.2	28.3	11.6	-16.7	72	.90	—
Qwen3.5-35B	0.6	35.8	28.9	-6.9	36	.81	—
<i>Human baseline (N=50)</i>							
Human	—	—	30.0	—	100	.72	—

(b) Single-gold illusion				
Model	Mode	Intended (true Acc)	Default (single-gold)	Δ
GPT-5	A	0.6	2.3	+1.7
	A+S	6.4	10.4	+4.0
	A+S+I	20.2	12.1	-8.1
GPT-4o	A	0.6	0.6	0.0
	A+S	12.1	37.0	+24.9
	A+S+I	4.6	6.4	+1.8
GPT-5-mini	A	0.6	0.0	-0.6
	A+S	0.0	0.0	0.0
	A+S+I	0.0	0.0	0.0

(c) Per-category targeting (AC@1)				
Model	Ent.	Met.	Rec.	Tmp.
GPT-5	.89	.92	.00	1.00
GPT-4o	1.00	.97	.00	1.00
GPT-5-m	.99	1.00	.00	1.00
Q3-30B	.93	.96	.00	1.00
Q3.5-35B	.85	.87	.00	1.00

and ignoring the ASC 606 practice of recognizing revenue at a point in time versus over time. This suggests that the domain-specific practice might not align too well with existing model’s generalized behavior, hence suggesting the need of specific tuning. We attach details and further case studies in in Appendix B. **Finding 3: Agents are not well-trained to notice ambiguity on financial-specific policies.**

4.5 Ambiguity category as training signal

While observing that the diagnosis is descriptive, we ask whether we can leverage the ambiguity category as training signal. We conduct conditioned training on the model, including fine-tuning on different categories followed by reinforcement learning on a Qwen3-4B policy, teaching the model to interact and target the correct category. Within the held-out split, we seek that the signals already jump significantly at supervised stage (Table 7, Figure 3a), and a full multi-turn reinforcement training further confirms that the episode is trainable end to end, with the clarifying turns rising from 5.7 to 8.6 (Figure 3b). The gain concentrates on the data-abundant categories, so the signal is actionable but still blinded in recognition policy in Sec. 4.4. Additional details are included in Appendix C.7. **Finding 4: FININTERACT is a trainable signal to address financial ambiguity.**

Table 7: RLVR ladder on FININTERACT using Qwen3-4B with held-out $n=51$ showing that we can train policy to clarify ambiguity using leak-proof AC@1 and interaction.

Stage	AC@1	Interaction	Accuracy	Reward
Base	25.0	54.9	84.3	0.96
+SFT	68.6	100	88.2	1.28
+KTO	70.6	100	84.3	1.27
+GRPO	66.7	100	88.2	1.30

5 Discussion

In our study, we observe that **categories are different learning problems**. While some clarification abilities are solved by scale, others, notably recognition basis, resist it, so a single aggregate score hides where the difficulty sits and curricula should budget data per category. On the other hand, **elicitation is a trainable skill current models lack**. Progress will come less from larger backbones, which barely move interaction accuracy, than from policies and objectives that reward asking on the right category at the right time. Nevertheless, even by including annotation for domain experts, it is still possible to retain **ambiguity in Human-AI collaboration** because models confidently return the default reading, motivating a safe behavior of asking instead of answering, hence the development of metrics that credit well-timed clarification and calibration process. While the current FININTERACT




Figure 3: Training on `FinINTERACT`. (a) Using AC@1 and interaction as training signal already improves performance in SFT. (b) True multi-turn GRPO on the 4B policy: mean reward and clarifying turns rise together over training.

is semi-automated, ideally a benchmark should be able to extend itself on categories and paired default/intended construction whose evidence is abundant, so as to **transfer to further categories, languages and other high-stakes domains** where ambiguity is commonly observed.

We note the following limitations of this study: **1) Scale** where a 173 instances corpus is considerably small given the need of construction rigor, **2) Language balance** where current corpus split between English and Chinese is 53 : 120, which could benefit models pre-trained in certain language, **3) Dominating categories** with entity scope ($n=94$) and metric definition ($n=63$), compared to the blind spot recognition policy with only $n=9$, which may suggest existing model is not well-trained enough to even generate such instance in the first place, **4) Generator and Evaluator bias** due to API budget constraints, as we hire GPT-5 family mainly for construction and simulation, which may bias towards the performance of GPT-5 family. **5) Data contamination** whereas instances are freshly constructed from FY2022-2025 facts where recently released models might have been pre-trained on. We include extended discussion in Appendix F.

6 Related Work

We include the most relevant works to `FinINTERACT` and defer additional financial and general-domain benchmarks as well as more widespread related work in Appendix D.

Financial QA benchmarks such as FinanceBench (Islam et al. 2023), FinQA (Chen et al. 2021), and DocFinQA (Reddy et al. 2024), together with recent agentic successors, enforce single-gold, single-turn evaluation and test numerical reasoning or retrieval without exploring whether a query is under-specified. To our knowledge, `FinINTERACT` is the first to make query ambiguity the central focus through paired default/intended interpretations and multi-turn interaction in the financial domain. We include comparison with existing financial QA benchmarks in Appendix D.

Ambiguous QA and clarifying questions. In the general domain, AmbigQA (Min et al. 2020), CLAMBER (Zhang et al. 2024), and QuestBench (Li, Kim, and Wang 2025) study

whether a model identifies under-specification and asks the minimal clarifying question, building on a long clarifying-question line in retrieval and dialogue. `INTERACTCOMP` (Deng et al. 2025) is our direct methodological ancestor, which we extend to finance with a domain-specific ambiguity taxonomy, per-category targeting (AC@1), and programmatically verifiable XBRL-grounded evidence pairs.

Multi-turn interaction and selective answering. Large models often *lose track* over multi-turn conversations (Laban et al. 2026), and calibrated systems frequently do better by answering *selectively* or abstaining under ambiguity (Cole et al. 2023; Kuhn, Gal, and Farquhar 2022). Our findings echo this in the financial setting, where forcing extra clarification degrades accuracy, but `FinINTERACT` turns the observation into a controlled, per-category diagnosis of *which* clarification an agent should issue and *whether* issuing it resolves the query, instead of only measuring whether more turns help.

7 Conclusion

We introduced `FinINTERACT`, the first bilingual benchmark for evaluating financial search agents on genuinely ambiguous queries. The five-category financial ambiguity taxonomy, paired default/intended interpretation design, and construction pipeline with adversarial verifier yield 173 high-quality instances to resist trivial disambiguation. Our central finding is that the key bottleneck of existing agent’s ability to handle financial QA ambiguity is interaction capability, but not financial knowledge. This capability contains two parts, i.e. the ability to ask for the right clarification and to update internal belief of the original answer. Specifically, when the ambiguity is fully solved, agents can answer with far higher accuracy at 93-95% accuracy. This amplifies the stagnation identified in `INTERACTCOMP` in the general domain, where the more terminology- and domain-specific financial industry requires specialized benchmark. We further turn this diagnosis into a training intervention in Sec. 4.5 to showcase that `FinINTERACT` also serves as an informative training signal, and release `FinINTERACT` with full construction, evaluation, and training code for reproducible follow-on work.

References

- Aliannejadi, M.; Kiseleva, J.; Chuklin, A.; Dalton, J.; and Burtsev, M. 2020. ConvAI3: Generating Clarifying Questions for Open-Domain Dialogue Systems (ClariQ). In *Search-Oriented Conversational AI (SCAI) Workshop at EMNLP*.
- Bigeard, A.; Nashold, L.; Krishnan, R.; and Wu, S. 2025. Finance Agent Benchmark: Benchmarking LLMs on Real-world Financial Research Tasks. *arXiv preprint arXiv:2508.00828*.
- Chen, Y.; Wang, Y.; Zhu, S.; Yu, H.; Feng, T.; Zhang, M.; Patwary, M.; and You, J. 2025. Multi-Agent Evolve: LLM Self-Improve through Co-evolution. *arXiv preprint arXiv:2510.23595*.
- Chen, Z.; Chen, W.; Smiley, C.; Shah, S.; Borova, I.; Langdon, D.; Moussa, R.; Beane, M.; Huang, T.-H.; Routledge, B.; and Wang, W. Y. 2021. FinQA: A Dataset of Numerical Reasoning over Financial Data. In *Proceedings of EMNLP*.
- Choi, C.; Kwon, J.; Lopez-Lira, A.; Kim, C.; et al. 2025. FinAgentBench: A Benchmark Dataset for Agentic Retrieval in Financial Question Answering. *arXiv preprint arXiv:2508.14052*.
- Cole, J. R.; Zhang, M. J.; Gillick, D.; Eisenschlos, J. M.; Dhingra, B.; and Eisenstein, J. 2023. Selectively Answering Ambiguous Questions. In *Proceedings of EMNLP*.
- Deng, M.; Huang, L.; Fan, Y.; Kong, F.; Zhang, J.; Ren, F.; Bai, J.; Yang, F.; Miao, D.; Yu, Z.; Wu, Y.; Zhang, Y.; Teng, F.; Wan, Y.; Hu, S.; Li, Y.; Jin, X.; Hu, C.; Li, H.; Fu, Q.; Zhong, T.; Wang, X.; Tang, X.; Tang, N.; Wu, C.; and Luo, Y. 2025. InteractComp: Evaluating Search Agents with Ambiguous Queries. *arXiv preprint arXiv:2510.24668*.
- Guo, X.; Xia, H.; Liu, Z.; Cao, H.; Yang, Z.; Liu, Z.; Wang, S.; Niu, J.; Wang, C.; Wang, Y.; et al. 2023. FinEval: A Chinese Financial Domain Knowledge Evaluation Benchmark for Large Language Models. *arXiv preprint arXiv:2308.09975*.
- Gwet, K. L. 2008. Computing inter-rater reliability and its variance in the presence of high agreement. *British Journal of Mathematical and Statistical Psychology*, 61(1): 29–48.
- Islam, P.; Kannappan, A.; Kiela, D.; Qian, R.; Scherrer, N.; and Vidgen, B. 2023. FinanceBench: A New Benchmark for Financial Question Answering. In *arXiv:2311.11944*.
- Krumdick, M.; Lorber, J.; Tsui, C.; Vithala, V.; Bhatt, M.; Kuo, F.; Madabushi, H. T.; and Mazumder, R. 2024. BizBench: A Quantitative Reasoning Benchmark for Business and Finance. In *Proceedings of ACL*.
- Kuhn, L.; Gal, Y.; and Farquhar, S. 2022. CLAM: Selective Clarification for Ambiguous Questions. In *ICML Workshop on Knowledge Retrieval and Language Models*.
- Laban, P.; Hayashi, H.; Zhou, Y.; and Neville, J. 2026. LLMs Get Lost In Multi-Turn Conversation. In *The Fourteenth International Conference on Learning Representations*.
- Li, B. Z.; Kim, B.; and Wang, Z. 2025. QuestBench: Can LLMs Ask the Right Question to Acquire Information in Reasoning Tasks? In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*.
- Li, Z.; et al. 2025. CondAmbigQA: A Benchmark for Conditional Ambiguous Question Answering. In *Proceedings of EMNLP*.
- Mateega, S.; Georgescu, C.; and Tang, D. 2025. FinanceQA: A Benchmark for Evaluating Financial Analysis Capabilities of Large Language Models. *arXiv preprint arXiv:2501.18062*.
- Min, S.; Michael, J.; Hajishirzi, H.; and Zettlemoyer, L. 2020. AmbigQA: Answering Ambiguous Open-domain Questions. In *Proceedings of EMNLP*.
- Rao, S.; and Daumé III, H. 2018. Learning to Ask Good Questions: Ranking Clarification Questions using Neural Expected Value of Perfect Information. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics (ACL)*.
- Reddy, V.; Balaguer, R.; Bhatt, G.; Bhatt, M.; Jiang, B.; Kuo, F.; Lee, C.-Y.; Madabushi, H. T.; Mazumder, R.; Mihindukulasooriya, N.; Morales, P.; and Picco, G. 2024. DocFinQA: A Long-Context Financial Reasoning Dataset. In *Proceedings of ACL*.
- Stelmakh, I.; Liang, Y.; Dziri, N.; Peng, N.; and Chang, K.-W. 2022. ASQA: Factoid Questions Meet Long-Form Answers. In *Proceedings of EMNLP*.
- Wei, J.; et al. 2025. BrowseComp: A Simple yet Challenging Benchmark for Browsing Agents. *arXiv:2504.12516*.
- Xie, Q.; Han, W.; Zhang, Z.; Lai, R.; Peng, M.; Lopez-Lira, A.; and Huang, J. 2024. FinBen: A Holistic Financial Benchmark for Large Language Models. In *Proceedings of NeurIPS Datasets and Benchmarks Track*.
- Yao, S.; Zhao, J.; Yu, D.; Du, N.; Shafran, I.; Narasimhan, K.; and Cao, Y. 2023. ReAct: Synergizing Reasoning and Acting in Language Models. In *Proceedings of ICLR*.
- Zamani, H.; Trippas, J. R.; Dalton, J.; and Radlinski, F. 2023. Conversational Information Seeking. *Foundations and Trends in Information Retrieval*.
- Zhang, T.; et al. 2024. CLAMBER: A Benchmark of Identifying and Clarifying Ambiguous Information Needs in Large Language Models. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics (ACL)*.

A Metric Details

This appendix expands the metrics defined in §3.5.

Accuracy grading. Accuracy is the primary ranking metric, mirroring INTERACTCOMP, BrowseComp, and FinanceBench. A GPT-4o-mini grader marks a final answer correct against the intended interpretation under financial-domain tolerance: a $\pm 1\%$ numeric tolerance, entity/ticker equivalence, and currency normalization, with a fiscal-year mismatch counted as wrong.

Interaction targeting (AC@1). For each interact action, a GPT-4o-mini classifier decides whether the clarifying question targets a true ambiguity category for that instance. We report the targeting accuracy AC@1 alongside the any-category, wrong-category, and generic-ask rates. Because a single clarification is often compound (one question can

name an entity, a period, and a metric at once), AC@1 has two forms: a lenient *touch* form (the question names any true category) and a strict *central* form (the single most-central category is a true category). The touch form is validated against human annotators (§3.4). These diagnostics characterize *how* a model interacts and do not rank models.

Calibration. The Expected Calibration Error (ECE) between stated confidence and correctness uses five equal-width confidence bins.

Clarification efficiency (DisE⁺). As a single supplementary number folding accuracy and clarification cost together, we report the efficiency of *successful* clarification,

$$\text{DisE}^+ = \mathbf{1}[\text{correct}] \times \mathbf{1}[n_{\text{asks}} > 0] \times \frac{H_0}{n_{\text{asks}}},$$

which is positive only when the agent asked at least one question and then answered correctly, and decreases with redundant asks. We restrict it to the asked-and-correct case to avoid the perverse incentive of an all-instances variant, which would award its maximum to a confident *zero-ask* answer, rewarding the very behavior the benchmark penalizes. DisE⁺ is a diagnostic of clarification efficiency, never the ranking metric.

B Detailed Results for the Core Findings

This appendix gives the full tables, figures, and findings summarized in §4.3, §4.2, and §4.4.

B.1 The Single-Gold Illusion

Before analyzing model behavior, we establish the core motivation empirically: *a conventional single-gold financial-QA benchmark cannot see the failure FININTERACT measures*. Existing benchmarks (FinanceBench, FinQA, DocFinQA) attach one gold answer per question. Built from the same filings, such a benchmark would almost always adopt the *default* reading, the headline figure a non-expert assumes, as that gold. Our paired construction lets us quantify the consequence directly: every instance carries both an intended and a default answer, so we re-grade the *identical* model outputs under both conventions (Table 6).

Finding 1: A single-gold convention over-states competence by up to 3×, and the gap is the benchmark’s reason to exist. In the retrieval-augmented regime where current single-gold benchmarks operate, the larger models return the *default* reading more often than the intended one: GPT-4o is graded 37.0% correct against the default but only 12.1% against the intended answer, a single-gold benchmark would report 3.1× its true competence. The model is not reasoning to the right answer; it is confidently producing a well-grounded answer to the *wrong reading of the question*, which a single-gold benchmark rewards and therefore cannot detect. (GPT-5-mini is the instructive exception: it captures *neither* reading because it asks rather than commits, so single-gold grading does not flatter it, but nor does it answer.) Critically, interaction is the *only* mode where intended accuracy exceeds default-capture (GPT-5: 20.2 vs. 12.1): asking is what converts a default-anchored guess into




Figure 4: **The elicitation gap.** Free-interaction accuracy (orange) versus the context-oracle ceiling (teal) per model. Once the ambiguity is resolved every model answers 90–95% correctly, so the 63–95 point gap is the cost of the model failing to *elicit* the disambiguation itself, not missing knowledge.

the correct intended answer. This is precisely the capability single-gold benchmarks are structurally blind to, and what FININTERACT is built to measure.

Human validation of the default. The single-gold argument assumes a conventional benchmark builder, seeing only the question and passage, would record the *default* reading as gold. We test this directly: five annotators, shown *Q* and the passage but not the resolving context *C*, independently record the answer they would treat as gold. They converge strongly (93% of items unanimous; mean inter-annotator agreement 0.98) and select the *default* reading 66% of the time versus the intended reading 34%, with fewer than 8% reporting that they noticed a second defensible answer. A single-gold benchmark built from these filings would therefore encode the default, empirically confirming the illusion’s premise rather than assuming it.

B.2 Latent Capacity and the Oracle Ceiling

To locate the bottleneck we bracket interactive accuracy with two ablations (Table 5 and Figure 4). **Forced interaction** requires the agent to ask $n=4$ clarifying questions before answering. **Context-oracle** hands the model the disambiguating context *C* (which fixes the intended interpretation) together with a passage containing *both* the intended and the default evidence spans, then asks it to answer, the ambiguity is pre-resolved but the model must still ground the correct value.

Finding 2: Models can answer once the ambiguity is resolved, they cannot resolve it themselves. The context-oracle ceiling is 93–95% for all three models, including GPT-5-mini, which scores 0% under free interaction. Latent answering capacity is therefore near-ceiling and roughly scale-invariant; what varies, and what FININTERACT measures, is whether a model can *resolve* the ambiguity end to end, eliciting the disambiguating context and then integrating it into a committed answer. Crucially, *forcing* clarification does not recover this capacity: requiring four asks leaves GPT-4o and




Figure 5: **Per-category clarification-capability profile** (touch AC@1 from Table 6). Every model saturates entity, metric, and temporal targeting but collapses to zero on `recognition_policy`, a blind spot shared across the closed frontier and the strongest open model.

GPT-5-mini near 0% (0.0% and 0.6%) and *lowers* GPT-5 from 20.2% to 1.5%, because forced asks consume the round budget on low-value questions without converging on a committed answer. This already points past pure elicitation: models frequently ask on-category yet still fail to commit the right value (the dominant error mode, §C.7), so the bottleneck is targeted, self-initiated disambiguation *and* the integration of the resulting answer, which neither scale nor brute-force asking supplies.

B.3 Per-Category Skills

The five categories are not just data strata; each names a distinct *skill*, the ability to recognize that a question is under-specified *along that category* and ask the question that resolves it. We make this measurable by decomposing each model’s per-category +*interact* competence into three components: **Recognition** (interaction rate, does it ask when it should?), **Targeting** (AC@1, does it ask about the *right* category?), and **Resolution** (accuracy, does asking yield the intended answer?). A model “has” a category skill only when all three are high; a low value localizes the failure to not-asking, asking-wrong, or asking-right-but-still-wrong. Table 6 reports Targeting, the decisive component, and Figure 5 shows the same profile as a per-category bar chart.

Finding 3: Ambiguity clarification is a per-category skill with a sharp touch versus central-targeting dissociation. Under the corrected classifier the pattern is precise. `recognition_policy` is a robust blind spot. No model ever asks about it, in either the touch or the central form (AC@1 = 0), so models essentially never question whether a figure is recognized point-in-time or over time. `temporal_scope` is fully solved. Models both touch it and make it the central concern (central AC@1 = 1.0). Entity and metric occupy a revealing middle. Models almost always *touch* them (touch AC@1 0.85–1.0) yet rarely make them the *central* point of the question (central AC@1 ≤ 0.16), defaulting instead to period confirmation. For these

two categories the failure is not neglect but imprecision, the model names the entity or metric in passing while centering the question elsewhere. Resolution follows the same shape across all five models (GPT-5 resolves 57% of temporal but $\leq 23\%$ of entity and metric, Figure 6). The gap is not uniform incompetence but missing specific targeting, which is precisely the intervention target (§C.7).

Per-model, per-category case analysis. Placing targeting and resolution side by side for every model and category (Figure 6) surfaces three patterns. First, resolution does *not* track targeting: GPT-4o and GPT-5-mini touch the correct category on entity and metric almost as often as any model (AC@1 0.94–1.0) yet resolve those categories at 0–5%, so a high AC@1 is no guarantee of a correct answer. Second, the model ranking *reorders* across categories. Among the full-scale models Qwen3.5-35B-A3B leads on the two abundant categories (entity 40%, metric 32%) while GPT-5 leads on temporal (57%). Third, the seven proprietary models resolve entity and metric substantially better than the older OpenAI ladder (entity 21–48%, metric 26–42%, versus GPT-5’s 23 and 13%), yet they inherit the *same* recognition blind spot, with touch AC@1 0 on every recognition instance they saw. `filing_vintage` is untestable everywhere ($n=0$). A single aggregate interaction score therefore conceals which category a given model can actually resolve, which is why we report per category.

C Additional Findings and Analyses

The main text reports the three headline findings (the single-gold illusion, the elicitation–integration ceiling, and the per-category skill dissociation). This appendix collects the remaining empirical findings and the category-conditioned interventions.

C.1 The Illusion Is Not Finance-Specific: A General-Domain Replication

Finding 4: the single-gold illusion is domain-general. Running the identical default-vs-intended measurement on AmbigQA (Min et al. 2020) (Table 8) reproduces all three effects outside finance, *consistently across three models*. Grading the *same* answer-only outputs against any accepted reading rather than a specified target overstates competence by 2.1–2.3 $\times$ (gpt-4o-mini 2.3, gpt-4o 2.2, gpt-5-mini 2.1), and the effect is robust to how the target reading is chosen, a random-target assignment gives 2.1 $\times$, versus our most-specific-reading default. Models return a valid but non-target reading 55–66% of the time, direct evidence the error is ambiguity-blindness, not noise. Elicitation is again under-used: even permitted to ask, models clarify on at most 27% of ambiguous queries and usually far less (gpt-4o 2.5%, gpt-5-mini <1%), and interaction leaves intended accuracy far below the any-valid ceiling (gpt-4o-mini 43.5 vs. 81). The one honest difference is the oracle ceiling, 50–69% here versus 93–95% in finance, because open-domain factoids stay hard even once disambiguated, so general knowledge is itself a limiter. This is exactly why finance is the sharper instrument: its exact, injectable ground truth isolates *elicitation* from *knowledge* in a way general-domain QA cannot. The

Full scale (n=173, top) | Proprietary pilot (n=50, bottom)

	(a) Targeting: AC@1 (touch)					(b) Resolution: +interact accuracy				
GPT-5	0.89	0.92	0.00	1.00	n/a	23	13	11	57	n/a
GPT-4o	1.00	0.97	0.00	1.00	n/a	5	2	11	14	n/a
GPT-5-mini	0.99	1.00	0.00	1.00	n/a	0	0	0	0	n/a
Qwen3-30B	0.93	0.96	0.00	1.00	n/a	12	10	11	14	n/a
Qwen3.5-35B	0.85	0.87	0.00	1.00	n/a	40	32	0	43	n/a
Claude-Sonnet-5	1.00	1.00	n/a	n/a	n/a	48	42	n/a	n/a	n/a
Gemini-3.5	1.00	0.94	n/a	n/a	n/a	28	37	n/a	n/a	n/a
Grok-4.5	0.83	1.00	n/a	n/a	n/a	34	32	n/a	n/a	n/a
DeepSeek-V4	0.96	1.00	n/a	n/a	n/a	21	32	n/a	n/a	n/a
GLM-5.2	1.00	1.00	n/a	n/a	n/a	41	42	n/a	n/a	n/a
Kimi-K3	1.00	1.00	n/a	n/a	n/a	41	26	n/a	n/a	n/a
GPT-5.6	1.00	1.00	n/a	n/a	n/a	34	26	n/a	n/a	n/a
	Entity	Metric	Recog.	Temp.	Filing	Entity	Metric	Recog.	Temp.	Filing

Figure 6: **Per-model, per-category performance.** (a) Targeting (AC@1 touch) versus (b) Resolution (+interact accuracy) for every model and ambiguity category. The top five rows are the full-scale panel ($n=173$, all categories) and the bottom seven are the proprietary pilot ($n=50$), for which only entity and metric carry enough instances ($n=29$ and 19), with recognition, temporal, and filing at $n \leq 1$ and shown n/a. `filing_vintage` is empty at full scale too ($n=0$). Every model targets entity and metric almost perfectly, the proprietary models resolve them better than the older OpenAI ladder, and no model targets recognition policy.

illusion itself, however, is a property of single-gold grading, not of finance.

C.2 Findings

We report results for the full OpenAI tier ladder, GPT-5, GPT-4o, and GPT-5-mini, across all three modes (Table 6). To probe whether the effect is an OpenAI-specific artifact, we additionally evaluate a preliminary seven-model, six-vendor panel of current frontier models (Claude-Sonnet-5, Gemini-3.5-flash, Grok-4.5, DeepSeek-V4, GLM-5.2, Kimi-K3, GPT-5.6) through a unified OpenRouter harness with an identical simulator/grader pipeline ($N=50/\text{mode}$, Table 6, panel (b)); it reproduces the headline pattern: added interaction is flat-or-negative for six of seven vendors and yields a clear gain only for Gemini, the weakest autonomous searcher. All headline accuracy comparisons below are reported with bootstrap 95% CIs over instances, and we test each within-model mode transition with a two-sided paired bootstrap ($n=173$). Six findings emerge here (Findings 5–10); the closely related latent-capacity ceiling is reported as a headline result in the main text (§4.2).

Finding 5: Interaction is a capability, not a free lever, it helps the strong model and harms the weak one. Adding interaction raises GPT-5 from 6.4% to 20.2% (95% CI [14.5, 26.6]; paired bootstrap $p<0.001$), but *lowers* GPT-4o from 12.1% to 4.6% ($p=0.002$). Forcing the weaker model to interact makes it significantly worse than not interacting at all. This extends INTERACTCOMP’s “interaction stagnation” into a sharper claim: for models that cannot use it, interac-

tion is actively harmful. Across all twelve evaluated models the interaction lever (added-interaction minus plain-search accuracy) is positive for only GPT-5 and Gemini and flat-or-negative for the other ten (Figure 7).

Finding 6: Model rankings flip with the interaction category. GPT-4o *outperforms* GPT-5 under pure retrieval (12.1 vs. 6.4) yet is *dominated* by it under interaction (4.6 vs. 20.2). A single-mode benchmark would mis-rank these models; the interaction category is what separates them, which is precisely the dimension FinInteract is designed to measure.

Finding 7: Retrieval cannot substitute for interaction. Oracle retrieval lifts GPT-5 to only 6.4%: because each disambiguating passage contains *both* the intended and default values, retrieval surfaces the ambiguity but does not resolve it. Only a targeted clarification collapses the question to a unique answer, which interaction (20.2%) begins to do.

Finding 8: The weak model’s failure is clarification non-integration, not inability to ask. GPT-4o asks readily (IR 99%, 4.7 questions on average) but fails to *integrate* the user’s responses: 87% of its interactive errors are non-committal refusals (e.g. “please specify the exact period”) issued *after* the simulator has answered. It can recognize ambiguity and ask, but cannot close the loop and commit, a multi-turn integration gap, not a protocol artifact (only 2% produced no answer at all).

Finding 9: Models are severely overconfident on ambiguous queries. Expected Calibration Error ranges from 0.53 (GPT-5, interact) to 0.95 (GPT-4o, answer-only); models report 85–95% confidence while answering 0–20% cor-

Table 8: The single-gold illusion replicates on AmbigQA (general domain, $n=200$, three models, *same* grading methodology; the grader retries with backoff so no answer is silently dropped). *Intended*: accuracy against a specified target reading. *Any-valid*: accuracy against *any* accepted reading, what a lenient single-gold benchmark credits. *Def.-cap*: fraction producing a valid but non-target reading. *Ask*: fraction that asks when permitted. gpt-5-mini interact omitted (it asks $<1\%$). Modes: *A* answer-only, *A+O* answer with context oracle, *A+I* answer with interaction.

Model	Mode	Intended	Any-valid	Def.-cap	Ask
gpt 4o-mini	A	31.5	72.5	55.0	0
	A+O	49.5	68.5	31.0	0
	A+I	43.5	81.0	53.0	27.0
gpt 4o	A	39.0	84.0	65.5	0
	A+O	68.5	81.5	31.5	0
	A+I	43.0	86.0	64.0	2.5
gpt 5-mini	A	36.0	77.0	62.5	0
	A+O	62.0	73.5	27.5	0

rectly. The remaining failure is not one of category *detection*. Under a human-validated classifier (agreement 0.75, matching inter-annotator agreement 0.73), both models almost always ask a question that touches the correct ambiguity category (AC@1 0.86 for GPT-5 and 0.94 for GPT-4o). What they lack is precise *central* targeting, only 0.06 and 0.04 of first questions make the true category their single central concern, and, above all, integration (Finding 8). The bottleneck is targeting precision and commitment rather than category detection, which motivates the category-conditioned policy of §C.7.

Finding 10: A smaller model collapses into ill-formed clarification, it asks the most but commits the least. GPT-5-mini scores 0.6/0.0/0.0% across the three modes. The cause is not ambiguity-blindness but the opposite: it is the *most* eager to clarify (interact IR 96%, 4.6 asks/instance), yet it issues *open-ended* questions (“which fiscal year? please paste the 2023 and 2022 figures”) rather than the protocol’s yes/no questions. Manually inspecting all 173 interact trajectories, 63% terminate on an unanswered open-ended question and 35% exhaust the 10-round budget still asking; only 1 instance ever commits a value (also wrong), and just 1.3% were format-unparseable. Even in Answer+Search, where no interact channel exists, it asks instead of answering, capturing *neither* the intended nor the default reading (0.0/0.0%). This is an extreme form of Finding 8’s integration gap and a distinct failure mode: clarification without commitment. The collapse is not an artifact of the yes/no protocol: with a *free-form* simulator that answers its open-ended questions from C (§5), GPT-5-mini still resolves 0% across all 50 instances (zero wrong-to-correct conversions), so the failure is integrating a received clarification, not the interaction grammar.

Language and category breakdown. Splitting by language confirms the benchmark’s bilingual-difficulty hypothesis: Chinese instances are consistently harder than English




Figure 7: **Interaction is a capability, not a free lever.** Added-interaction minus plain-search accuracy across the twelve full-scale and cross-vendor models. Interaction improves accuracy for only GPT-5 and Gemini and is flat or negative for the other ten, so the ability to convert clarification into a correct answer is rare.

in the interact mode (GPT-5 18.3% ZH vs. 24.5% EN; GPT-4o 0.8% vs. 13.2%), reflecting more complex corporate-structure naming and metric conventions in A-share filings. Stratifying by primary category (interact mode), *metric_definition* is the hardest well-populated category (GPT-5 12.7%, $n=63$) and *entity_scope* the next (23.4%, $n=94$), while *temporal_scope* is comparatively easy (57.1%, $n=7$). This inverts the general-domain pattern where entity disambiguation dominates, and identifies metric-definition ambiguity as the distinctive financial challenge.

Difficulty and headroom. Per-instance accuracy falls monotonically with disambiguation entropy H_0 (from 10% at $H_0=1$ to under 1% at $H_0\approx 3.6$), validating H_0 as a difficulty proxy; **no full-scale model solves 79 of 173 instances (46%) under free interaction.** These are not broken items: 96% of them become solvable under the context-oracle, so they are hard specifically at the elicitation step, not invalid or unanswerable. With the best free-interaction accuracy at 28.9% full scale (and 44% in the cross-vendor pilot), the benchmark is far from saturated. The supplementary clarification-efficiency metric mirrors the accuracy ranking in interact mode ($\text{DisE}^+ = 0.125$ for GPT-5 vs. 0.039 for GPT-4o): GPT-5 not only answers more ambiguous queries correctly but does so with fewer wasted asks.

Table 9: Zero-shot *central*-targeting (central-form AC@1) by base-model scale (Qwen3, +*interact*, denominator = instances where the model asked). Under the validated touch form every scale already reaches ≈ 0.9 on entity and metric, so scale buys the precision to *center* the question on the right category rather than detection. Entity central-targeting scales steeply and is largely solved at 32B, while recognition policy stays 0 at every scale, including the same 32B that targets entity well. Recognition is underpowered ($n \leq 9$ asked-on) but flat throughout.

Category	4B	14B	32B
entity_scope	.00	.00	.68
metric_definition	.05	.00	.17
recognition_policy	.00	.00	.00

C.3 Human Baseline

To bound the task from above under *realistic* interaction, where the disambiguating context must be *elicited*, not handed over, we collect a human baseline on a stratified 50-instance subset (32 EN / 18 ZH, balanced by category). Two finance-literate annotators (one bilingual) face the same input a model does in the search regime: the ambiguous question and a retrieved passage containing both candidate values, position-shuffled and unlabeled. Each may ask *one* yes/no question, which the experimenter answers truthfully from the withheld context C (Yes/No/IDK), before committing a final answer. This mirrors the agent’s *search*→*interact*→*answer* loop, so the human row is directly comparable to the +Interact column; we grade it with the identical GPT-4o-mini grader.

Finding 11: Even humans who can ask resolve only 30%, the task is hard, not merely mis-specified for models. Annotators reach 30.0% accuracy overall (34.4% EN, 22.2% ZH), above every full-scale model (which top out at 28.9%) but far from the 93–95% context-oracle ceiling. The gap between the oracle ceiling and the human baseline is the cost of *elicitation*: knowing the resolved answer is near-trivial, but extracting the right bit through a single clarification is hard even for experts. Across models and the human baseline, targeting and resolution *dissociate*: the models with the highest AC@1 (GPT-4o and GPT-5-mini at 0.94) post the lowest accuracy, while the human, at a lower touch AC@1 of 0.72 (one focused question rather than a sweep of several categories), posts the highest (Figure 8). Notably the annotators never targeted the recognition-policy category either (AC@1 0 on its eight items), so the blind spot is not model-specific. The asymmetry is sharpest in Chinese, where annotators still capture the default reading 72.2% of the time at only 22.2% accuracy. Chinese instances are therefore harder to *resolve*, not harder to *recognize*, the *gui-mul-kou-fei* (parent-attributable vs. deduction-of-nonrecurring) net-income distinction is sticky even once correctly surfaced, corroborating the language-asymmetry caveat in §5 with a mechanism.

Asking on target does not imply resolving




Figure 8: **Targeting does not imply resolution.** Accuracy against AC@1 (touch) in the interact regime, for the five full-scale models and the human baseline (star). The relationship is inverted: models that ask on the correct category most often resolve the ambiguity least often, whereas the human asks less precisely yet resolves more. Asking the right question and answering it are separate abilities.

C.4 Skill Dissociation Across Model Scale

Finding 12: the categories are different kinds of skill, some scale-solvable, one a scale-invariant blind spot. A controlled scale ladder (Qwen3 4B → 32B, Table 9) separates the categories sharply, once we look at *central* targeting. *recognition_policy* is a scale-invariant blind spot. It is never targeted at any scale, in either the touch or the central form (0.00 at 4B, 14B, and 32B, including the same 32B that targets entity well), so more parameters buy no recognition-basis targeting at all. *entity_scope* and *metric_definition* tell a subtler story. Every scale already *touches* them (touch AC@1 ≈ 0.9), but the ability to make them the *central* point of the question emerges only with scale, entity central-targeting rising from 0.00 at 4B and 14B to 0.68 at 32B. Scale therefore does not buy ambiguity detection for these categories, which is already present, but the precision to center the question on the right category. The dissociation confirms the taxonomy carves genuine capability boundaries rather than arbitrary strata, and it pinpoints recognition basis as the category a targeted intervention must address (§C.7).

C.5 Qualitative Case Studies

Three instances make the failure modes concrete, one for each category where the pattern is sharpest. Every case pairs a natural analyst question with two defensible readings whose ground truth is exact and machine-verifiable.

Entity scope (resolvable, yet over-elicited). “What was Meta Platforms’ operating income?” is \$46.75B for the total consolidated company and \$62.87B for the Family of Apps

segment, a 34% gap on the same GAAP metric and fiscal year. A search-only agent retrieves a passage containing both figures and commits to one, and single-gold grading credits it whenever that figure happens to match the key. This is the category models handle best, and the base policy often resolves it directly (§C.7), which is precisely why forcing a clarifying question on entity scope can *lower* accuracy rather than raise it.

Recognition policy (the blind spot). “What were Tesla’s revenues?” is \$97.69B under GAAP accrual revenue recognized on the income statement, but only \$14.92B under a cash-basis proxy taken from operating cash flows, a 6.5× difference driven entirely by the recognition basis. This is the category models almost never surface: instead of asking which recognition basis is meant, the typical clarification reverts to the more salient period or entity question (§4.4), so the 6.5× ambiguity is never raised and the default is reported as fact. The category is represented in the network yet not acted on (§C.6).

Metric definition (clarification supplies the answer). “What was Workday’s operating income growth rate?” is 126.8% on a GAAP basis and 26% on a non-GAAP basis for the same year-over-year comparison. Here the base model cannot answer correctly without knowing which definition is intended, so a single on-category question converts a wrong answer into a right one, the mechanism behind the +45 point training-time gain on this category (§C.7).

Together the three cases trace the arc of the paper: a category that is resolvable but over-elicited (entity), a category that is represented yet never acted on (recognition), and a category where eliciting the definition is the whole battle (metric).

C.6 Mechanistic Case Study: Where Ambiguity Lives in the Network

The behavioral results (§4) show that frontier models default to the naive interpretation rather than clarifying. A natural question is whether models internally *represent* this ambiguity at all. Because this requires access to hidden states, we probe four open models as a proxy, spanning two size classes and two architectures: **Qwen3-4B-Instruct** (dense, 37 layers), **Qwen3.5-4B** (hybrid linear/full attention, 33 layers), and the mixture-of-experts **Qwen3-30B-A3B** (49 layers) and **Qwen3.5-35B-A3B** (41 layers).

Method. For each instance we form a within-item contrast: the bare ambiguous question Q versus Q concatenated with its disambiguating context C . We define a per-layer *ambiguity direction* as the difference in mean final-token activations between the two conditions, and score any hidden state by its (mean-centered) projection onto that direction; tracing this across layers gives a depth profile of how strongly each layer separates an under-specified question from a disambiguated one. Independently, on the two larger models we train a per-layer linear probe to decode the ambiguity *type* (entity_scope vs. metric_definition), which, unlike a context-presence detector, cannot be solved trivially.

Result. Figure 9(a) overlays the depth profiles (normalized for differing depth and scale). On all four models the signal

is near-zero through most of the stack and rises sharply to a peak in the *final few layers* ($\approx 95\text{--}98\%$ relative depth: layers 35/37, 32/33, 48/49, 40/41). Figure 9(b) shows that the ambiguity *type* is linearly decodable at 0.80 accuracy (majority baseline 0.60) on both larger models, though it becomes available at different depths, early ($\approx 27\%$) in Qwen3-30B and late ($\approx 85\%$) in Qwen3.5-35B. Across four models, two scales, and three architecture families, the representation is consistent: financial-query ambiguity, and its category, is linearly encoded, emerging toward the top of the network.

Representation, elicitation, and the cost of interacting.

The two MoE models we probe are also *evaluated* (Table 6): beyond linearly encoding the category, both *elicit*, and they differ in *how* they ask. Qwen3-30B interacts on 72% of instances but loops without committing in Chinese (76% of its interactive transcripts never produce a final answer), whereas Qwen3.5-35B interacts less (36%) but commits in 95% of cases and targets the gold category more often (AC@1 0.40 vs. 0.16). With the full retrieval corpus in place their accuracy is now measured, and it exposes a sharper pattern: plain search reaches 28.3%/35.8%, but *adding* interaction *lowers* accuracy to 11.6%/28.9% ($-16.8/-6.9$ points). Interaction is thus net-negative relative to search for both open models, as it is for GPT-4o (12.1 $\rightarrow$ 4.6), and only the strongest model, GPT-5, converts clarification into a gain (6.4 $\rightarrow$ 20.2). An interpretation-oracle control (the resolved interpretation supplied but no evidence) scores $\approx 0\%$ for both MoE models, so the residual gap to the oracle ceiling is an evidence-grounding cost, not a parametric-recall one. The representational and behavioural evidence shows the category is both registered and acted on at the level of asking. A linear-probe replication on the *base* (non-instruct) checkpoints reproduces the decodability and the early-vs-late depth split (peaks 0.71/0.74 vs. baseline 0.60).

Finding 13: Open models linearly represent both ambiguity and its category, replicating from 4B to 35B and across dense and MoE architectures. The representation is type-specific (entity vs. metric decodable at 0.80) and robust to scale and architecture, yet behaviorally these same models still default rather than clarify. The behavioral gap is thus better read as a failure to *act* on a represented ambiguity than a failure to register it, though establishing this causally (e.g. via activation steering) remains future work.

Scope and caveats. This is mechanistic evidence on open proxies, not the closed frontier models benchmarked in §4; diff-in-means probing and linear decoding establish *availability*, not causal use; and category-type decoding is supported only for the two well-populated categories (entity_scope, metric_definition). We deliberately do *not* draw behavioral conclusions from these open models’ free-form generations, where surface cues (chain-of-thought markers, retrieval that surfaces the evidence span) confound naive clarification counts; the asking-behavior claims come solely from the graded closed-model evaluation in §4. A further caveat is that the ambiguity direction is a difference between Q and $Q+C$ inputs, which also differ in length and surface lexicon, so it may partly encode the presence of added context rather




Figure 9: Four-model depth profile of the ambiguity signal (Qwen3 4B→35B, dense and MoE). (a) Ambiguous—disambiguated separation rises to a peak in the final layers of every model (dotted verticals mark per-model peaks; normalized for depth and magnitude). (b) The ambiguity *type* (entity vs. metric) is linearly decodable at 0.80 on both MoE models, well above the 0.60 majority baseline, peaking at different depths.

than ambiguity as such. Matched neutral-context controls, a bag-of-words lexical control, a cross-company split, and causal steering are needed before reading the probe as more than evidence that category information is *linearly decodable* from the representation. That the category is *represented but not acted on* frames the bottleneck as an action problem on an available signal, and motivates the category-conditioned interventions of §C.7, which target exactly that signal.

C.7 Diagnosis and Intervention: Acting on the Ambiguity Category

Our results converge on a single diagnosis. The benchmark localizes the bottleneck to *eliciting* the disambiguation: agents answer 93–95% correctly once the intended interpretation is supplied (§4.2) but far less on their own (a mean of 8% across the OpenAI ladder), and the mechanistic probe shows why, open models *linearly represent* the ambiguity *category* yet do not act on it (§C.6). The remaining gap is thus an *action* problem on an already-available signal, and the *category* is its natural unit. We now turn that diagnosis into intervention: we first catalog how acting fails (a seven-type error taxonomy), then present two interventions that share one idea, **condition on the category**, instantiated at inference time (Category-Aware ReAct, §C.7) and at training time (category-guided SFT/GRPO, §C.7).

Error Taxonomy Beyond aggregate accuracy, we characterize model failures using a seven-type error taxonomy (Table 10) derived from FININTERACT’s evaluation signals. Each type is diagnosable from existing metrics without additional annotation.

Diagnostic Mapping from Evaluation Signals The existing evaluation metrics map directly onto error types, enabling failure attribution without additional annotation:

- **Low IR + low Acc** → E1 (ambiguity blindness dominant)

Table 10: Error taxonomy E1–E7 with diagnostic signals and improvement directions.

Type	Description	Improvement direction
E1	Ambiguity blindness: no interaction, wrong answer	Ambiguity pre-check; train on detection
E2	Wrong-category clarification: asked wrong dimension	Category classifier; financial taxonomy grounding
E3	Generic clarification: vague ask, no category target	Category-conditioned clarification templates
E4	Retrieval dominance: searched but skipped interaction	Force interpretation check after retrieval
E5	Clarification misuse: right question, wrong answer	State-gated answering; stronger grounding
E6	Evidence grounding: template-oracle correct, agent wrong	Improved source-grounded extraction
E7	Over-interaction: correct but low DisE ⁺	Penalize DisE ⁺ /AC@1 via efficiency signal

- **High IR + low AC@1** → E2/E3 (interacts but misdirected)
- **High AC@1 + low Acc** → E5/E6 (right category, wrong grounding)
- **High GenericAskRate** → E3 (lacks financial clarification structure)
- **High WrongCategoryRate** → E2 (confidently misdirected)
- **Always-ask** >> **Standard ReAct** → E1 dominant; latent capacity exists
- **Template-oracle** >> **Agent-interact** → “what to ask” is

the bottleneck

- **Category-oracle** $\gg$ **Agent-interact** $\rightarrow$ category recognition is the bottleneck
- **Template-oracle still low** $\rightarrow$ E6 residual; evidence extraction still fails

Misintegration, not misdirection, dominates the error budget. Attributing each interact-mode outcome to the taxonomy (Figure 10) shows that *E5* (right category touched, wrong answer) accounts for the large majority of failures, 67% for GPT-5, 89% for GPT-4o, 90% for GPT-5-mini, and 75% for Qwen3-30B, while not-asking (E1) and off-category asks (E2 and E3) together stay below 15% for the closed models. The bottleneck is converting a correctly targeted clarification into the intended answer, not deciding whether or what to ask. Qwen3.5-35B-A3B is the informative exception, asking less (E1 21%) but resolving more of what it attempts (Correct 35%), so its failures lean toward under-asking rather than misintegration. This corroborates Findings 3 and 8 across the full error budget and motivates the state-gated answering of §C.7.

Inference-Time Intervention: Category-Aware Clarification ReAct The first intervention conditions on the category *without changing the model weights*. Standard ReAct leaves the agent free to decide whether and what to ask. We propose **Category-Aware Clarification ReAct**, a three-component structured interaction policy that directly addresses E1–E3.

Component 1: Ambiguity Pre-Check. Before any action, the agent explicitly decides which of the five ambiguity categories are underspecified in the question and initializes an interpretation state $\mathcal{S} = \{\text{entity, period, metric, basis, vintage}\}$, marking each field as *resolved* or *unknown*.

Component 2: Category-Conditioned Clarification. For each *unknown* field, the agent generates a targeted yes/no question using category-conditioned templates (one per category) rather than free-form clarification. The agent emits at most one clarification per turn, prioritizing the most information-theoretically valuable unknown field.

Component 3: Interpretation-State-Gated Answering. The agent may only emit a final answer once all required fields are resolved (or the user explicitly declines to clarify). This prevents the *retrieval dominance* pattern (E4) where the agent answers immediately after finding a plausible number in a search result.

Expected outcomes. Relative to standard Answer+Search+Interact, Category-Aware ReAct should reduce GenericAskRate (E3), WrongCategoryRate (E2), and IR=0 failures (E1), while improving AC@1 and ultimately Accuracy. Template-oracle provides an upper bound: if Category-Aware ReAct approaches template-oracle performance, it confirms that category recognition is the dominant remaining bottleneck after the E1–E3 errors are addressed.

Pilot results. We evaluate the policy on a GPT-5 backbone over a stratified $N=50$ subset run through the same OpenRouter harness, alongside four matched baselines (Standard

ReAct, Always-Ask, Category-Oracle, Template-Oracle). Category-Aware ReAct reaches 46.0% accuracy against the 34.0% of Standard ReAct in the same run, a +12 point gain, and posts the best calibration of any condition (ECE 0.241 versus 0.308 for Standard ReAct). Always-Ask collapses to 10.0%, confirming that asking indiscriminately is worse than not asking and that the gain comes from asking *selectively* on the right category. The single-category diagnostics (AC@1, WrongCategoryRate) are not informative for this policy: because the Ambiguity Pre-Check deliberately enumerates every candidate category to rule each in or out, an any-category touch form saturates near 1.0 while the strict central-category form is near 0, so neither cleanly credits the policy and we omit AC@1 for it in Table 6. The accuracy and calibration gains are the load-bearing result. We report this as a pilot rather than a full-panel row because it is a single backbone at $N=50$, and we expect the effect size to tighten at full scale. One caveat bears directly on the mechanism: the policy bundles all three components above, and we do not ablate them here, so part of the +12 may come from the state-gated commit (Component 3, an integration fix) rather than from better category-conditioned elicitation (Component 2). A generic-structured clarification baseline that keeps the structure but drops the named categories, together with a per-component ablation, is the clean control and remains future work.

Training-Time Intervention: Category-Guided SFT and GRPO The per-category skill analysis (§4.4) gives this intervention a precise target: models do not lack clarification ability uniformly, they lack *specific* category skills (AC@1 = 0 on recognition policy, partial on metric). The question is whether those missing skills can be installed by training rather than prompting. The second intervention applies the *same* idea, condition on the category, but in the weights rather than the prompt, showing the benchmark is not only diagnosable but *trainable*. Starting from Qwen3-4B-Instruct (4-bit QLoRA), we run an RLVR ladder, supervised finetuning (SFT), KTO, then GRPO, on the FinInteract training split, with all user simulators, graders, and category judges served locally (zero API cost). We score each stage with the two *leak-proof* metrics, AC@1 and interaction rate (Table 7); accuracy is reported only with the caveat that the training environment’s retrieval returns the disambiguated evidence span, so it is partly leaked and is *not* a reliable learning signal.

The category-guided teacher. The decisive ingredient is how SFT demonstrations are generated. An uninformed teacher, even a 72B model, almost always asks the *obvious* ambiguity (typically the reporting period) and rarely the subtle gold category (entity scope, recognition policy), yielding $\sim 0\%$ on-category demonstrations. Privately revealing the gold category to the teacher while generating demonstrations (the hint is *not* stored in the trajectory) lifts the on-category yield to $\sim 79\%$, a reusable recipe for any category-structured clarification task.

SFT does the heavy lifting. The leak-proof signals jump almost entirely at the supervised stage, AC@1 rises $25 \rightarrow 69$ and interaction $55 \rightarrow 100$ (Figure 3a), while KTO and GRPO move within run-to-run noise at $n=51$. Training teaches the model to *always interact and target the correct category*, the




Figure 10: **Error-type distribution** in `+interact`, attributing every instance to the taxonomy of Table 10 with the corrected category classifier. *E5 misintegration* (the agent asks, touches the correct category, yet answers wrong) dominates for every model, while ambiguity blindness (E1), wrong-category (E2), and generic (E3) asks stay comparatively rare.

behavior `FININTERACT` is built to elicit. Whether this gain reflects a sharpened internal *representation* of the category or only a rewired output *policy*, probing the trained model as in §C.6, is a direct test of the diagnosis and a target for future work.

True multi-turn GRPO. The ladder above optimizes a single-turn approximation. We additionally confirm the full multi-turn episode is trainable end-to-end: optimizing the 4B policy over entire ReAct episodes (loss masked on simulator and retrieval tokens) for 15 steps raises mean reward $0.55 \rightarrow 1.08$ while clarifying turns rise $5.7 \rightarrow 8.6$ (Figure 3b), the category-aware reward is teaching the policy to ask more. We report this as a learning-signal demonstration, not a converged model.

Caveats. Results are on a held-out split of $n=51$; all judges and the teacher are local Qwen2.5-32B/72B models (weaker than GPT-4o, and siblings of the policy); the terminal reward is partly leaked; and GRPO is run for tens of steps, not to convergence. We therefore frame this as evidence that `FININTERACT` is *trainable*, a category-guided SFT seed instills category-targeted clarification, rather than as a state-of-the-art accuracy result.

When Can the Benchmark Refresh Itself? Three Co-Evolution Probes A benchmark that models eventually memorize stops measuring. Because `FININTERACT` couples a grounded data-generation pipeline (§3.3) to a trainable policy (§C.7), it can in principle *co-evolve*: generate fresh instances on a category the current model fails, train on them, and repeat, a data $\leftrightarrow$ model loop in the spirit of self-improving generators (Chen et al. 2025), but *grounded* in real filings rather than the ungrounded self-play those methods risk. We run three gated rounds, on the scarce, subordinate

category (`recognition_policy`), the abundant, salient category (`entity_scope`), and the abundant-but-only-weakly-scale-learnable category (`metric_definition`, §4.4), to ask not *whether* the loop works but *when*.

All three rounds run end-to-end: the Constructor $\rightarrow$ Verifier pipeline yields hard instances (every one passes the adversarial gate: the base fails it answer-only) and category-guided SFT converges. The outcome splits cleanly by *abundance*, not by scale-learnability. On both abundant categories, `entity_scope` and `metric_definition`, Targeting generalizes to held-out *human* items ($.00 \rightarrow .68$ and $.00 \rightarrow .58$, each $\sim 4\times$ the $+15$ bar) and to held-out *generated* items, with no catastrophic forgetting. On the scarce `recognition_policy` it does not move ($.00$ on both held-out sets): the trained model reverts to the salient period/entity question rather than the recognition-basis distinction. The contrast rules out a broken pipeline, the *same* machinery installs entity and metric but not recognition, and, because metric is only *weakly* scale-learnable yet trains as well as entity, it rules out scale-learnability as the gate (§C.7).

Finding 14: the loop closes on abundant categories and fails on scarce ones; whether elicitation helps is a separate, per-category property. A self-refreshing `FININTERACT` is viable wherever the data arm can supply enough frontier instances, both abundant categories install and generalize; the scarce category (recognition’s 36-instance ceiling from just 30/1,847 primary passages) does not reliably (Finding 17). But installing on-category elicitation is not automatically beneficial: it *helps* metric (accuracy $15 \rightarrow 60$, $+45$: the base cannot answer without the metric definition, so the clarifying question *supplies* it: 18 wrong $\rightarrow$ correct, 0 regressions) and *hurts* entity (accuracy $32.5 \rightarrow 22.5$, -10 : the base can

Table 11: Three co-evolution rounds (category-guided SFT on a 4B policy). Targeting (AC@1) is on a *held-out human* probe never seen in training. The loop closes on both *abundant* categories and fails only on the *scarce* one; whether asking *helps accuracy* flips on whether the base could already answer.

	entity abundant	metric abundant	recognition scarce
Frontier instances generable	99	76	36 (ceiling)
Targeting, base $\rightarrow$ evolved	.00 $\rightarrow$.68	.00 $\rightarrow$.58	.00 $\rightarrow$.00
on held-out <i>generated</i>	.40	.88	.00
Interaction rate	40 $\rightarrow$ 100	50 $\rightarrow$ 100	67 $\rightarrow$ 100
De-leaked accuracy	32.5 $\rightarrow$ 22.5	15 $\rightarrow$ 60	11 $\rightarrow$ 11

often answer the entity directly, so forcing a question turns six correct direct answers into wrong multi-turn ones). Both drove interaction to 100%; the accuracy sign flips on whether the base could already answer, the training-time face of the “interaction can hurt versus plain answering” effect seen at inference (§4.2). The actionable step, which we demonstrate next, is a *conditional when-to-ask gate*, elicit only where the model cannot already answer, plus multi-round curriculum and fresh-filing ingestion to lift the scarce categories above their availability ceiling.

Finding 15: the over-elicitation regression is fixable at inference, but not by a train-time gate. We test the when-to-ask gate directly on the frozen human probes ($n=40$ per category). A *train-time* gate that mixes no-ask demonstrations into the category-guided SFT backfires: even 7 of 164 no-ask demonstrations tip the policy into a bimodal collapse in which it answers directly without retrieving, which is fatal because in FININTERACT almost nothing is answerable from parametric memory (entity accuracy 32.5 $\rightarrow$ 7.5, metric 60 $\rightarrow$ 2.5). An *inference-time* gate recovers the loss with no extra training. Routing by the base model’s own ask decision, taking the co-evolved asker’s answer when the base chose to ask and the base’s direct answer otherwise, lifts entity back to 30.0 from the always-ask 22.5 (matching the base’s own 32.5) at a 40% ask rate rather than 100%, and keeps +20 of metric’s +45 gain (35.0 versus base 15.0). Over-elicitation is thus an inference-time routing problem rather than a training one. The residual metric gap traces to the base’s uncertainty signal under-asking where metric needs a question, so a confidence-calibrated gate that asks when the silent answer would be *wrong* rather than when the base is merely unsure is the concrete next lever.

What Gates Co-Evolvability? A Falsification Test The three rounds let us isolate *which* category property predicts whether the loop closes. Two candidates are natural, a category’s *scale-learnability* (does zero-shot Targeting rise with model size?) and its *source-abundance* (can the data arm supply frontier instances?), and we had also hypothesized a representational version: *learnable iff the category is linearly decodable* from hidden states. The *metric_definition* round was designed to discriminate the first two; a linear probe tests the third. Both tidy hypotheses fail.

Scale-learnability does not gate co-evolution. *metric_definition* is abundant (498 primary passages) but only *weakly* scale-learnable (zero-shot AC@1

.05 $\rightarrow$.26, slope +0.21, versus entity’s +0.77). We pre-registered the two-factor prediction that it would co-evolve only *partially*. It did not: metric is a *full* GO (Targeting +0.55, matching entity’s +0.61; Table 11). What separates GO from NO-GO is not the scale-slope but *abundance*, both abundant categories close the loop; only the scarce category fails.

Representation does not gate it either, and is not the bottleneck. A per-layer one-vs-rest linear probe decodes every category from base hidden states, and the pattern *anti-tracks* learnability: *recognition_policy*, the unlearnable, NO-GO category, is the *most* decodable (AU-ROC 0.93 on a size-matched sample, above entity/metric at ~ 0.77); decodability is *flat* across the 4B $\rightarrow$ 32B sweep even as behavioural AC@1 climbs .12 $\rightarrow$.89; and category-guided SFT leaves it unchanged. The category is fully *represented* at every scale, before and after training: the model simply does not *act* on it, the “represents-but-doesn’t-act” gap of §C.6, now shown to hold even for the category co-evolution cannot fix.

Finding 16: co-evolvability is gated by source abundance, driven by salience, not by scale-learnability or representational decodability. What survives both falsifications is a single upstream cause. *recognition_policy* is a candidate category in 796 passages but *primary* in only 30 (salience 0.04): it is almost always the *subordinate* ambiguity, behind a more salient entity or period category in the same passage. Low salience produces both the scarcity that starves the data arm and the salient competitor the model’s prior reaches for instead, so the representation is present but never acted on and cannot be cheaply trained in. Salient categories behave oppositely (entity, salience 0.92: abundant and trainable). For a practitioner, **salience, not scale curves or probe accuracy, is the cheap, up-front signal** of which categories a co-evolving benchmark can keep hard and which need fresh human curation. The corollary for evaluation is sharper: the hardest ambiguity to make models *handle* is not the one they cannot represent, but the one that is always someone else’s more-obvious question.

Finding 17: salience predicts a category’s data-need, not just a binary outcome, and a validation-guarded arms race confirms it. Extending each category to a multi-round curriculum (Fig. 11), with checkpoints selected on a frozen validation set and reported on the held-out human probe, refines the binary gate into a gradient that tracks salience. **Entity** (salience 0.92) learns category-targeting from a *sin-*

Table 12: What predicts co-evolvability. The observed outcomes track *abundance/salience*, *not* the AC@1 scale-slope (metric: low slope, GO) and *not* decodability (recognition: highest AUROC, yet does not co-evolve). Salience = primary/any-candidate passages; decodability = peak one-vs-rest AUROC (recognition on a size-matched augmented sample; raw $n=9$ overfits to 1.0). *metric is a GO but needs ~ 2 accumulated slices (~ 50 instances); one 30-instance slice teaches it nothing (Finding 17). ‡ recognition is *unreliable* rather than a clean fail: held-out human AC@1 ranges 0–0.44 across training runs on the scarce $n=9$ probe.

Category	Primary	Salience	Slope	Decode	Co-evolve
entity_scope	905	0.92	+0.77	0.77	GO
metric_definition	498	0.48	+0.21	0.77	GO*
temporal_scope	150	0.15	sat.	0.92	solved
filing_vintage	264	0.33	–	–	untestable
recognition_policy	30	0.04	+0.00	0.93	weak ‡




Figure 11: Multi-round co-evolution with a validation guard. *Left*: validation AC@1 (selection metric) per round. *Right*: held-out *human-probe* AC@1 (never used for selection). Entity learns targeting from a single 30-instance slice and transfers to humans (0.63 $\rightarrow$ 0.93); metric needs *accumulation* (one slice teaches nothing, brown, flat, while two accumulated slices reach 0.75, green); recognition stays at base (≈ 0.05) even with a fair 26-instance set. The guard preserves human transfer (it rolls back the round-3 candidate, which regresses on validation).

gle 30-instance slice and transfers to human ambiguity at 0.88–0.93. **Metric** (0.48) needs *accumulation*: one slice yields 0.0, two accumulated slices reach 0.75, so its earlier single-round GO used ~ 60 instances precisely because that is roughly what the category requires. **Recognition** (0.04) is *unreliable*: across training runs its human AC@1 ranges only 0–0.44 on the scarce $n=9$ probe, weak and inconsistent, not a clean zero. Scoring the same runs as an *Elo arms race* between the solver and a hard-example miner (win = AC@1, which, unlike recall-walled resolution, varies across checkpoints) makes the instrument non-degenerate: entity is *solver-runaway* (solver Elo 1000 $\rightarrow$ 1551, outpacing a corpus-bounded miner that cannot synthesize adversarially harder targeting), while recognition *stalls* (flat). Two guarded rounds peak; the third overfits and is rolled back. The practical reading: salience is an up-front estimate not just of *whether* an category co-evolves but of *how much* data it needs, and one well-targeted round, guarded on held-out validation, is close to the sweet spot.

D Extended Related Work

This appendix expands the related work summarized in §6.

Financial QA Benchmarks. Table 13 compares FININTERACT to major financial QA benchmarks. FinanceBench (Islam et al. 2023), FinQA (Chen et al. 2021), DocFinQA (Reddy et al. 2024), BizBench (Krumdick et al. 2024), FinBen (Xie et al. 2024), and FinanceQA (Mateega, Georgescu, and Tang 2025) all enforce single-gold-answer conventions and evaluate models in a single turn; FinanceQA further shows that frontier models fail $\sim 60\%$ of realistic analyst tasks, but attributes this to numerical rigor rather than ambiguity. FinEval (Guo et al. 2023), a large Chinese financial-knowledge benchmark, incidentally identifies an “ambiguity handling weakness” error type in its agent tasks but does not isolate or control it. Recent *agentic* financial benchmarks raise task realism, Finance Agent Benchmark (Bigear et al. 2025) (537 expert research tasks, where the best model reaches only 46.8%) and FinAgentBench (Choi et al. 2025)

Table 13: Comparison of financial QA benchmarks. ✓= supported; – = not applicable or not evaluated.

Benchmark	Scale	Ambig.	Interact.	Multi-turn	Biling.	Evidence	Grounding
FinanceBench (Islam et al. 2023)	150	–	–	–	–	human	SEC
FinQA (Chen et al. 2021)	8,281	–	–	–	–	human	10-K/Q
DocFinQA (Reddy et al. 2024)	7,437	–	–	–	–	human	10-K/Q
BizBench (Krumdick et al. 2024)	multi-task	–	–	–	–	auto	mixed
FinBen (Xie et al. 2024)	multi-task	partial	–	–	partial	auto	mixed
FinanceQA (Mateega, Georgescu, and Tang 2025)	multi-task	–	–	–	–	expert	SEC
Finance Agent Bench. (Bigeard et al. 2025)	537	–	–	–	–	expert	SEC
FinAgentBench (Choi et al. 2025)	26K	–	–	–	–	auto	SEC
FinEval (Guo et al. 2023)	8,351	–	–	–	ZH	expert	exams
FININTERACT (ours)	173	✓	✓	✓	EN+ZH	LLM+human	EDGAR+CNINFO

(agentic retrieval over EDGAR), but both remain single-turn and single-gold, evaluating analysis and retrieval rather than whether an agent recognizes that a query is under-specified. FININTERACT is the first to make query ambiguity the central, controlled construct, via paired default/intended interpretations and multi-turn interaction, in the financial domain.

Ambiguous QA. AmbigQA (Min et al. 2020), ASQA (Stelmakh et al. 2022), CondAmbigQA (Li et al. 2025), and CLAM (Kuhn, Gal, and Farquhar 2022) have matured in the general (Wikipedia) domain. More recent benchmarks target the agent’s *own* ability to ask: CLAMBER (Zhang et al. 2024) evaluates whether models identify and clarify ambiguous information needs, and QuestBench (Li, Kim, and Wang 2025) tests whether they can ask the minimal necessary question in reasoning tasks. Both are general-domain and judge a single clarifying step; FININTERACT adds *per-category* targeting (AC@1), programmatically verifiable financial grounding, and, via the co-evolution rounds (§C.7), an analysis of *when* asking helps versus hurts. We adapt this line’s formalism to finance, replacing human-judged ambiguity labels with programmatically verifiable XBRL-grounded evidence pairs.

Clarifying Questions. Asking clarifying questions is a long-standing goal in information retrieval and dialogue. Rao and Daumé III (Rao and Daumé III 2018) rank clarification questions by expected value of perfect information; ClariQ (Aliannejadi et al. 2020) established a shared task for generating clarifying questions in open-domain dialogue; and the mixed-initiative paradigm is surveyed in Conversational Information Seeking (Zamani et al. 2023). FININTERACT instantiates this lineage with a leak-proof, per-category measure of *which* clarification an agent should issue and a verifiable outcome for whether issuing it resolves the query.

Interactive Search Agents. INTERACTCOMP (Deng et al. 2025) is the direct methodological ancestor. We extend its general-domain setting to finance with a domain-specific ambiguity taxonomy and bilingual evaluation. BrowseComp (Wei et al. 2025) tracks longitudinal retrieval trends; we test whether the interaction-capability gap identified there persists in finance.

E Extended Discussion

This appendix expands the discussion points summarized in §5.

Elicitation is a trainable skill that current models lack. Our central result is that models fail to resolve ambiguity not for want of knowledge but for want of asking the right question (RQ2). This reframes the problem for future work. Progress will come less from larger backbones, which barely move interaction accuracy, and more from policies and objectives that reward asking on the correct category at the right time. The Category-Aware ReAct policy and category-guided fine-tuning (RQ3) are first steps, and both leave a large gap to the oracle ceiling, so reward shaping on the ambiguity category and calibrated deferral are natural next directions.

Categories are different learning problems. The per-category dissociation shows that some clarification abilities are solved by scale while others, notably recognition basis, resist it and are only unreliably teachable. A single aggregate interaction score therefore hides where the real difficulty sits. Future benchmarks and training curricula should treat each category as its own learning problem and budget data accordingly, which our salience estimate begins to quantify.

Human and AI collaboration under ambiguity. Because models confidently return the default reading, the safe behavior is often to ask rather than to answer. This motivates

evaluation and interfaces that credit well-timed clarification and calibrated abstention rather than raw commitment, and it connects to mixed-initiative retrieval where the cost of a confident wrong answer is high.

Self-refreshing benchmarks and broader scope. Our co-evolution probes suggest a benchmark can extend itself on categories whose evidence is abundant, guarded by held-out validation, which offers a path to keep the instrument fresh as models improve. Natural extensions include more ambiguity categories, additional languages and filing regimes, and transfer of the paired default and intended construction to other high-stakes domains where a single gold answer hides genuine ambiguity.

F Extended Limitations

This appendix expands each limitation summarized in §5.

Scale. At 173 instances FinInteract is small (comparable to FinanceBench’s 150 (Islam et al. 2023)); we prioritize per-instance construction rigor (every item is adversarially verified and R14-discriminating) over raw size, and Appendix G shows the pipeline scales at \$0.31/instance.

Language balance. The English split (53) is smaller than the Chinese split (120); per-language results (§4) should be read with this asymmetry in mind, and the EN portion is a target for growth.

Category coverage and recognition-policy validity. The taxonomy is intentionally five-category, but `filing_vintage` is unrepresented and `temporal_scope/recognition_policy` are sparse (Table 4): viable structured ambiguity for these categories is rare in public filings (Appendix G). We therefore treat entity scope ($n=94$) and metric definition ($n=63$) as the two well-powered categories and the other three as exploratory. An item-level audit of the nine `recognition_policy` instances is a further caution: most are ASC 606 timing disaggregations (revenue recognized point-in-time versus over-time), where neither value is the total a non-expert would assume, and one pairs total GAAP revenue against an operating cash-flow proxy that is not a defensible reading of “revenue.” Even our human annotators never targeted the category (AC@1 0 on its items). We therefore report the recognition-policy blind spot as a preliminary observation on a small, imperfect sample, not a headline result, and recommend rebuilding this category with analyst-audited pairs before drawing quantitative conclusions.

Model coverage. We report the OpenAI tier ladder (3 models) and two open MoE models at full scale ($n=173$); a seven-model, six-vendor panel of current frontier models is additionally evaluated in a preliminary pilot ($N=50$ /mode, Table 6, panel (b)), which reproduces the central finding but at reduced statistical resolution, and a full-scale multi-vendor evaluation remains in progress. The human baseline ($N=50$, a single yes/no clarification) is a constrained reference rather than an upper bound, since models given up to ten rounds can exceed it, and it is collected from two annotators on a subset. A larger expert panel remains future work.

Grader and simulator. Correctness and category-targeting are judged by GPT-4o-mini and the user is sim-

ulated by GPT-5. We validate these against humans on a stratified, blindly double-annotated sample (five annotators). The correctness grader agrees with human consensus at Cohen $\kappa=0.85$ (raw agreement 0.92), exceeding inter-annotator agreement (0.86), so accuracy figures hold under human labels. The AC@1 classifier initially agreed only weakly (0.53). A single-label design over-assigned `temporal_scope` to compound questions that also name an entity or metric. The corrected multi-label *touch* classifier reaches 0.75 agreement, matching inter-annotator agreement (0.73), and the AC@1 values reported here use it. Because this shows models mostly do ask on-category, the corrected reading is that the gap lies in central targeting and integration rather than detection. A strict central-category variant is discriminative but not yet human-validated, so we report it only as a supplementary decomposition. Interaction results are also robust to the simulator: on a 50-instance subset, standard +Interact accuracy is 34/28/26% under the LLM, deterministic, and 15%-noisy simulators. Crucially, the yes/no action grammar does not manufacture the weak-model collapse: with a *free-form* simulator that answers open-ended clarifying questions from *C* (for example “which fiscal year?”) rather than only yes/no, GPT-5-mini stays at 0% across all 50 instances (identical to the yes/no simulator, zero wrong-to-correct conversions, three-round budget). Its failure is therefore a post-clarification integration failure, not an artifact of the interaction protocol.

Construction circularity and synthetic intent. The constructor, adversarial verifier, and user simulator are all GPT-5-family models, and the verifier removes instances the GPT-5 family answers correctly without *C*. This depresses GPT-5-family answer-only accuracy (0.6%) by construction relative to the non-OpenAI vendors the filter does not target (mean 10.0% answer-only across six vendors), so cross-vendor answer-only comparisons are not fully apples-to-apples and the GPT-5-family floor is a filter-induced lower bound. The elicitation gap itself is unaffected: the filter targets answer-only solvability, not interactive resolution, and every vendor still resolves far below its oracle ceiling. Verifying a subsample with a non-OpenAI adversarial ensemble is a natural robustness check for future work. A related construct-validity caveat is that questions are reverse-constructed from filings rather than drawn from real user queries, and the “intended” interpretation is a constructor-assigned hidden context. The benchmark therefore measures a model’s ability to recover a *specified* interpretation, a necessary and controllable proxy for, but not identical to, disambiguating the intent of a real analyst.

Contamination. Instances are freshly constructed from FY2022–2025 structured facts and answer-only accuracy is $\approx 0.6\%$ across all sources, but we do not claim absolute contamination safety for the DocFinQA-derived minority.

All instances are ambiguous by construction. Every one of the 173 items is built to require clarification, so a model can succeed by assuming ambiguity is always present. FinInteract therefore measures *what* to ask and *how* to integrate the answer, not a calibrated *when*-to-ask policy: it does not score false-positive clarification, over-asking on already-specified questions, or the correct choice to answer directly when a

mild ambiguity does not change the value. Adding matched unambiguous and answer-invariant controls is the single most important extension for turning the benchmark into a test of calibrated interaction, and we accordingly frame our claims around asking-and-integrating rather than when-to-ask.

Intervention evidence is preliminary and not component-attributed. The Category-Aware ReAct gain (34 $\rightarrow$ 46%) is a single GPT-5 backbone on $N=50$ (about six additional items), without a significance test, and its matched-baseline subset scores above the full-scale GPT-5 row, so the effect must be confirmed at full scale with confidence intervals. Just as important, the policy bundles three components (ambiguity pre-check, category-conditioned questioning, and interpretation-state-gated answering) that we do not ablate: the improvement may stem substantially from the state-gated commit, an *integration* fix, rather than from better category *elicitation*, so the pilot does not by itself isolate the taxonomy’s mechanistic role. A generic-structured clarification baseline and a per-component ablation are the needed controls. The training-time results are likewise early, resting on a small held-out probe, gains dominated by supervised fine-tuning with only a short reinforcement phase, an optimistic oracle-retrieval reference, and interaction driven to near-total compliance rather than a learned when-to-ask policy.

A single protocol confounds ability with interface adherence. All vendors share one ReAct prompt and yes/no action grammar. GPT-5-mini’s 0% is partly a protocol-adherence failure (it keeps asking open-ended questions, exhausts the round budget, and rarely commits a final answer) rather than proof of zero financial-ambiguity ability. A free-form simulator rules out the yes/no grammar as the cause (the collapse persists), but tool schema, few-shot examples, constrained decoding, and model-specific adapters all shape protocol compliance, so a fixed cross-vendor prompt is a conservative rather than a maximally fair comparison; per-model interface tuning is future work.

Scope. The paper spans a benchmark, the single-gold illusion, the taxonomy, and several exploratory studies (representation probing, inference and training interventions, co-evolution, and salience). The three main-text findings are the load-bearing contribution and the rest are explicitly exploratory and appendix-scoped; the two highest-value additions they point to, matched ambiguous/unambiguous evaluation and a component ablation of the intervention, are our priority follow-ups.

G Curation Cost to Replicate the Benchmark

A practical concern for any LLM-assisted benchmark is the cost to reproduce or scale it. We instrument the construction pipeline (§3.3) and report measured API costs. All figures use list prices at construction time (per 1M tokens: GPT-5 \$1.25 in / \$10.00 out; GPT-5-mini \$0.25 / \$2.00; GPT-4o-mini \$0.15 / \$0.60) and count API fees only, they exclude the one-time filing-collection compute and the human spot-check labor (10–20% of accepted instances).

The constructor, not the verifier, is the yield bottleneck. Re-running the pipeline over a held-out sample of candidate

passages reproduces a sharp funnel (Table 14). Only 11.5% of keyword-tagged candidate passages survive the constructor, which rejects the rest as containing *no viable ambiguity* for the target category; but of those that survive, 98.5% pass the 10-trial adversarial verifier. In other words, the construction prompt is strict and the verifier rarely needs to fire, the cost of the benchmark is dominated by constructor calls spent on candidate passages that turn out not to support a genuine ambiguity, not by verification.

Table 14: Measured curation cost. Unit costs are from instrumented pipeline runs (constructor-only $n=20$; full pipeline $n=6$); the acceptance funnel is from the full construction run ($n=1127$ candidate passages). A full pipeline pass is one GPT-5 constructor call plus ten verifier trials (five GPT-5, five GPT-5-mini) with GPT-4o-mini grading.

Quantity	Value
Constructor viability (viable / candidate)	11.5%
Verifier acceptance (accepted / viable)	98.5%
Overall acceptance (accepted / candidate)	11.3%
Cost per rejected candidate (constructor only)	\$0.020
Cost per viable passage (full pipeline)	\$0.155
Blended cost per candidate passage	\$0.035
Cost per accepted instance	\$0.31

Projected cost to scale. At the measured blended rate, producing N accepted instances requires $\approx N/0.113$ candidate passages and the following API spend: 173 instances (the current release) $\approx$ \$54; 500 instances $\approx$ \$155; 1000 instances $\approx$ \$309. The benchmark is therefore inexpensive to scale in API terms, an order of \$300 takes it to 1000 instances, and the true constraint is the supply of candidate passages that contain structured, verifiable ambiguity. This favors *better passage pre-selection* (raising the 11.5% viability rate) over brute-force compute, and explains why the rare categories (e.g. *filing_vintage*), whose candidate passages almost never survive the constructor, are the expensive ones to grow.

H Construction and Validation Details

This appendix expands the data sources, curation pipeline, and validation summarized in §3.

Data sources. **English (EN, $\sim 80\%$ of pool).** *DocFinQA*, 780 long-context passages extracted from 10-K and 10-Q filings, with verified candidate answers from annotated QA pairs. *EDGAR (XBRL-derived)*, 111 passages built from machine-readable XBRL facts for S&P 500 and Russell 1000 companies (FY2024–2025). **Chinese (ZH, $\sim 20\%$ of pool).** *CNINFO/akshare*, 237 passages derived from A-share annual reports (SSE 50 + CSI 300) via structured financial-statement APIs, in three passage types: (1) metric definition, *kou-fei* vs. *gui-mu* (non-recurring-excluded profit vs. parent-attributable profit); (2) entity scope, *gui-mu* vs. *jing-lirun* (parent-attributable vs. fully consolidated including minority interest); (3) temporal scope, year-over-year revenue growth. ZH instances are expected to be harder for current models due to sparser training data on A-share filings and more

complex corporate-structure nomenclature (e.g., A/H-share distinctions, *jituan* (group) vs. *zi-gongsi* (subsidiary) scope).

Three-role construction pipeline. Step 1: Constructor (GPT-5). A single GPT-5 call generates the full (Q, C, A) triple given a filing passage, verified candidate answer, and mandatory primary category. The output JSON includes question, context, answer, default_answer, intended_evidence_span, default_evidence_span, intended_interpretation, default_interpretation, and axes_exercised, with category-specific instructions injected per primary category. **Step 2: Pre-verifier sanity checks.** Fourteen automated code-level rules (Table 15) are applied without API calls, and any failure immediately rejects the passage. **Step 3: Adversarial verifier.** Ten verifier calls run in parallel (five GPT-5-mini and five GPT-5); each answers Q without C and states its assumed interpretation, and an instance is rejected if ≥ 2 of the 10 both produce A and align with the intended interpretation. This two-condition criterion prevents false rejections from lucky guesses.

Category diversity enforcement. We enforce target category shares through four mechanisms: (1) a priority function with a $3\times$ penalty on over-quota categories; (2) hard exclusion at $2\times$ quota; (3) a rarity-adjusted initial sort (passages sorted by $\sum_{\text{categories}} (\text{target_share}/\text{pool_frequency})$); and (4) a dynamic re-sort every 10 instances.

Human validation protocol. Two finance-literate annotators (one bilingual, for the Chinese split) independently review instances on an eight-question protocol: H1 *Is Q ambiguous without C ?* H2 *Is the default interpretation plausible for a non-expert?* H3 *Does C uniquely identify the intended interpretation?* H4 *Does C avoid directly stating the answer?* H5 *Is A correct under the intended interpretation?* H6 *Is A_d correct under the default interpretation?* H7 *Is the primary category label correct?* H8 *What yes/no question would a financial analyst naturally ask to resolve the ambiguity?* (free text, used to compute human AC@1). Acceptance requires H1–H6 to hold after adjudication by the primary author. On a stratified 60-instance sample (by language and category), per-item raw agreement is 0.82–0.97. Because the labels are heavily skewed toward “yes” ($> 85\%$ per item), Cohen’s κ is degenerate under the Feinstein–Cicchetti prevalence paradox, so we report Gwet’s AC1 (0.85 average, range 0.78–0.97; 1.00 for H7). The single consensus-rejected instance (consensus rejection 1.7%) was a category over-reach, a revenue-recognition-timing pairing whose plausible default is the firm’s total revenue rather than a disaggregation component. It is answerable only under context-oracle, where removing it shifts the ceiling by $< 0.05\text{pp}$ (93–95% unchanged).

I Construction Quality-Control Rules

The pre-verifier applies fourteen automated code-level rules before any API call (Table 15), rejecting a candidate the moment any rule fails.

Table 15: Pre-verifier sanity checks (14 rules).

Rule	Description
R1	Q must not be answerable yes/no
R2	Q must not contain disambiguating terms or inline dates
R3	A taken verbatim from the source candidate value (not invented)
R4	Q must name the company
R5	Q must ask about a substantive financial metric
R6	Answer type consistent with question type
R7	Context C must not contain the answer value
R8	Intended and default interpretations must differ
R9	All category labels must be valid vocabulary items
R10	Answer must be non-trivial (not empty or “0”)
R11	default_answer present and differs from answer
R12	Both evidence spans non-empty
R13	Evidence spans must be meaningfully distinct ($\leq 85\%$ overlap)
R14	Intended and default answers differ beyond grader tolerance

J Reproducibility

Models and decoding. Closed models are accessed via their official APIs (evaluated 2026): gpt-5, gpt-4o, gpt-5-mini. Reasoning models (gpt-5 family) use default temperature with a generous max_completion_tokens headroom (+4096 above the output budget) so internal reasoning does not truncate the answer; other models use temperature 0 for the agent and grader. **User simulator.** GPT-5 at temperature 1.0, constrained to {Yes, No, I don’t know} and conditioned on the disambiguating context C . **Grader.** GPT-4o-mini at temperature 0, binary correctness with finance tolerance ($\pm 1\%$ numeric, entity/ticker equivalence, currency normalization; fiscal-year mismatch = wrong). **Protocol.** ReAct with search/interact/answer actions, max 10 rounds; the forced-interaction ablation requires $n=4$ asks before answering; the context-oracle ceiling supplies C plus both evidence spans (position-shuffled). **Statistics.** Accuracy CIs are 95% bootstrap over instances ($B=2000$); mode transitions use a two-sided paired bootstrap ($B=5000$). All evaluation, construction, and analysis scripts are released.

K Datasheet and Licensing

Provenance. English instances derive from public SEC EDGAR filings (10-K, FY2022–2025) and a small DocFinQA-sourced set; Chinese instances derive from public CSRC annual reports retrieved via CNINFO/akshare. All source documents are public regulatory filings; no proprietary, paywalled, or personal data is used. **Schema.** Each instance is a JSON record with fields: instance_id, language, source, ticker, company, filing_type, filing_date, question, context (C), answer, default_answer, intended_evidence_span, default_evidence_span,




Figure 12: **FININTERACT distributions.** (a) Primary ambiguity category: the corpus is entity/metric-dominant (54% / 36%), with the three subordinate categories sparse, mirroring how often each category is the *primary* ambiguity in public filings (§2). (b) Difficulty by disambiguation entropy $H_0 = \log_2 k$: instances span one to 3.58 bits, most at one or two bits with a substantial multi-category tail at 3.58 bits (45 instances).

intended_interpretation/default_interpretation
 ({entity, period, metric, basis}), axes, n_axes, h0, and
 qc (per-rule pass flags + verifier blind-solve rate). **Intended
 use.** Evaluating whether interactive search agents recognize
 and resolve query ambiguity; not a training corpus and not
 financial advice. **License and release.** The benchmark,
 construction pipeline, and evaluation code are released for
 research use; redistribution of source filings follows the re-
 spective public-domain (EDGAR) and regulatory-disclosure
 (CSRC) terms. **Maintenance.** The frozen v1 release is
 versioned; corrections and the planned EN-split expansion
 will be issued as numbered releases.

Disclosure of Generative AI Use

Generative AI tools were used in this work in two distinct roles. First, as *components of the research artifact itself*: large language models act as the adversarial constructor, the user simulator, and the automated grader described in Section 3 and Section 4, and every such use is documented in the paper and the accompanying code with the specific model, version, and prompt. Second, generative AI tools were used to assist with *manuscript preparation*, limited to language editing, L^AT_EX formatting, and code refactoring. All research questions, experimental design, analyses, and scientific claims are the authors’ own, and the authors have independently verified every result, table, figure, and citation reported here. The authors take full responsibility for the entire content of this paper, in accordance with the AAAI Code of Professional Ethics and Conduct.